

Ember+ Specification

Version: 2.40
Revision: 11
Date: 2016-10-04
Authors: Marius Keuck, Philip Boger

Document History

2.40 – 10/2016

- Updated the content table.

2.40 – 05/2016

- Changed the reference to the ASN.1 encoding rules documentation.

2.40 – 05/2016

- Added the “Number” field to the Template type.

2.40 – 04/2016

- Added chapter about the Template extensions (Ember+ 1.4, Glow 2.40).

2.31 – 02/2016

- Added a more detailed description about the online state of a node and the expected behaviour of a provider and consumer.

2.31 – 03/2015

- Added a more detailed comment about the format string.

2.31 – 02/2015

- Corrected an incorrect wording on page 59.

2.31 – 04/2014

- Added support for multiple schema identifiers.

2.30 – 01/2014

- Added chapter about the schema extension and a section about node removal.

2.20 – Revision 4 – 11/2013

- Added more detailed information about reporting a node without children.

2.20 – Revision 3 – 10/2013

- Corrected information on function content

2.20 – Revision 2 – 10/2013

- Corrected information on node and parameter numbering

2.20 – Revision 1 – 05/2013

- Reviewed information on very simple providers and the GetDirectory command.

2.20 – 05/2013

- Added Chapter about Function Extensions (Ember+ 1.2, Glow 2.20)

2.10 – Revision 1 – 10/2012

- Added a small section about the three layers of Ember+ to the Introduction.

2.10 – 10/2012

- Added Chapter about Matrix Extensions (Ember+ 1.1, Glow 2.10)

2.5 – Revision 5 – 10/2012

- Added more details to chapter “Message Length” (pg. 41)

2.5 – Revision 4 – 09/2012

- Updated some code examples which still used the insertElement method instead of insert.

2.5 – Revision 3 – 08/2012

- Added the naming rules for the node and parameter identifier string (see [The identifier property](#)).
- Added the information that replying to a Keep-Alive request is mandatory (see [Keep-Alive mechanism](#)).
- Added the “isOnline” property to the Node

2.5 – Revision 2 – 08/2012

- Added a document revision number.
- Added a more detailed description for the Subscribe and Unsubscribe commands.

07/2012

- Added a note about the formula format. More details can be found in the document “Ember+ formulas”.
- Described a rule to detect empty nodes.

06/2012

- Updated the Glow DTD version to 2.5, which introduces a more flexible way to transport streams. Besides, the types QualifiedNode and QualifiedParameter have been added, which offer a more compact method to transmit property changes.

05/2012

- Changed the numbers of the Subscribe and Unsubscribe command.
- Updated the Glow DTD version to 2.4.

04/2012

- The tags of nodes and parameters are now context specific.
- A node can now be marked as root.
- Added the EnumMap property and the description of the EnumCollection and StringIntegerPair types.

02/2012

- Removed the tag information in the command list since a command is only represented by its number.

02/2012

- Initial draft.

Content

Introduction	8
EmBER.....	10
Introduction.....	10
About BER	10
Objectives	10
Conformance.....	10
References.....	11
Basics.....	11
Overview	11
Tagging [X.690 8.14]	11
Length forms [X.690 8.1.3]	11
Container Usage	11
Basic document structure.....	12
Types.....	12
Overview	12
Boolean [X.690 8.2]	12
Integer [X.690 8.3]	12
Real [X.690 8.5]	13
UTF8String [X.690 8.21]	13
Octet String [X.690 8.7]	13
Relative Object Identifier [X.690 8.20].....	13
Set [X.690 8.11]	13
Sequence [X.690 8.9, X.690 8.10]	13
Implementation recommendations.....	14
Encoder.....	14
Decoder.....	14
Glow specification	15
Introduction.....	15
Data types and properties	15
The tag format.....	15
Application defined tags	16
Application defined types.....	17

Glow specific properties	22
Application defined commands	29
Ember+ 1.1: Matrix Extensions	31
Introduction	31
Type Definitions	31
Use Cases	38
Performance Characteristics	43
Ember+ 1.2: Function Extensions	44
Introduction	44
Type Definitions	44
Use Cases	46
Ember+ 1.3: Schema Extensions	48
Purpose.....	48
Identification within an Ember+ Tree.....	48
Scope.....	48
Schema definition	48
Application Notes	49
Syntax of Schema Identifiers.....	49
References.....	50
Ember+ 1.4: Templates.....	51
Purpose.....	51
Specification within an Ember+ tree.....	51
Restrictions, semantics and precedence	52
GetDirectory requests.....	53
Compatibility and migration	53
Changes to the Ember+ DTD.....	54
Ember+ usage.....	56
Representing a device.....	56
Querying a data provider	57
The EmBER library	61
Compiling the library	61
Including the library	61
Library structure	61

The util namespace	62
The BER namespace.....	62
The Dom namespace	63
The Glow namespace.....	63
Using the glow classes.....	64
Creating a GetDirectory request.....	64
Replying to a GetDirectory request	64
Changing a parameter value	65
Reporting a parameter value change	65
Traversing a tree.....	66
Behaviour rules and other guidelines	68
General behaviour rules	68
Message length	68
The identifier property.....	68
The description property	68
The number property	69
When to use the element number or identifier	69
Keep-Alive mechanism.....	69
Number of consumers per provider	69
Notifications	69
Parameter value range changes	70
Value change requests	70
Remarks on the Behaviour of Embedded Devices	70
Removing Sub-Trees at Runtime.....	70
General recommendations.....	70
Separate parameters into logical categories.....	70
Avoid indexed parameters.....	71
Offline trees.....	71
Frequently asked questions	72
Why are the node and parameter numbers session based?.....	72
Message Framing	73
The S101 protocol.....	73
Encoding a message.....	73

Decoding a message.....	73
CRC Computation.....	74
S101 Messages.....	75
Message types	75
Commands.....	75
Messages.....	75
Glow DTD ASN.1 Notation	77
Appendix	93
S101 CRC Table	93
S101 Decoder Sample.....	93
S101 Encoder Sample	94

Introduction


The Ember+ protocol has been designed to allow the communication between two endpoints, one being the data provider and the other one being the consumer. The data provider usually is a piece of hardware which offers a set of controllable parameters, while the consumer may be a control- or monitoring-system which provides access to these parameters. This document describes the Ember+ protocol specification and the encoding it uses. Furthermore it provides some implementation hints and gives several code examples that demonstrate the use of this protocol.

The Ember+ standard defines three separate data representation technologies, which when used in conjunction produce a data layout that conforms to the rules defined in this document:

- **Glow** is the name of the data schema that defines the data types used to convey Ember+ information.
- **EmBER** is the name of the encoding used to store instances of the types defined by *Glow*.
- **S101** is the framing protocol used to transmit *EmBER* encoded data, ensuring message integrity.

When comparing the transmission of a well-formed Ember+ message to the transmission of a standard webpage, the following layers of data representation are equivalent:

- *Glow* and XHTML
- *EmBER* and XML
- *S101* and HTTP



Ember+ Layers

The chapter "[EmBER](#)" describes the basics of the *Embedded Basic Encoding Rules* (EmBER), which is the encoding used by this protocol. The code samples listed in this document refer to the EmBER library implementation written in C++. The library is freely available and usually published together with this document. It implements the subset of BER which is required for this protocol. BER is a well-known encoding standard. It is platform independent, generates small data packets and offers an XML-like flexibility. It provides encoding rules for all common data types, like integers, decimal numbers, strings, byte arrays and more. Additionally, it provides container types called set and sequence. EmBER uses a TLV (Type, Length and Value) system in order to specify what kind of data is currently present in the encoded data buffer. This allows extending the protocol by user defined data types, similar to XML where a node may have any name but also indicates what kind of attributes or data it contains.

The chapter "[Ember+ specification](#)" provides detailed information about the types, properties and commands being used. The Ember+ specification relies on BER and defines a set of objects in ASN.1 notation, which are also listed and described in this document. One of the important points Ember+ put a focus on is to keep communication between the peers as generic as possible. That means that the consumer doesn't have to have any kind of detailed knowledge of the device it is connected to. Instead, the device is responsible of providing its structure when the consumer requests it. This allows a consumer to control any kind of device as long as they follow the specification.

Chapter number three, called "[Ember+ Usage](#)" describes how a provider represents its data, how a consumer queries the structure and how it can change a parameter value.

In "[EmBER Library](#)" the different layers of the C++ library implementation are described. Additionally, the section shows how the library can be used to encode or decode data.

The chapter "[Message Framing](#)" describes the framing protocol which is used to transport encoded Ember+ packets, called S101. The library also contains a basic encoder and decoder framework that can be used to read and write S101 messages.

"[Behaviour rules and guidelines](#)" lists how a consumer and a provider must behave in several situations. Besides, this chapter contains some general guidelines about how a tree should look like.

The Ember+ object definition is called "Glow"; the specification is listed in the chapter "[Glow DTD ASN.1 Notation](#)".

The library and the framework referred to by this document provide everything needed in order to easily implement the Ember+ protocol. It makes it unnecessary to have detailed knowledge about the encoding and decoding algorithms.

Additionally to the C++ library mentioned and used in this document, there are also implementations written in standard C (called `libember_slim`) and a .NET version.

EmBER

Introduction

This chapter describes the data types supported in this library and the restrictions.

About BER

BER stands for Basic Encoding Rules. The rules describe how data is being encoded and decoded. The encoding algorithms are fast and produce small data packets, which is a good reason to use this encoding on embedded systems. Each data entity is encoded by a tag, the data length and the value itself, in short terms TLV. The tag simply defines the type, while the length field contains the length of the encoded value in bytes. Both, the tag and the length are encoded as well, which reduces the size of the resulting data packet even more. The structure of a tag is described in the second chapter.

Objectives

EmBER has been developed in order to provide a way to encode data for transfer between different electronic devices with varying hardware resources, independent from the transmission medium and protocol. Ember can also be used to serialize data to non-volatile memory (e.g. a file on a hard disk).

Ember has been designed to provide the following assets:

1. Conformance to a well-known and widely adopted standard
2. Compact size of the encoded data to minimize transmission load
3. Platform independence
4. XML-like flexibility
5. Binary storage of values for fast encoding and decoding

Conformance

Ember forms a subset of the Basic Encoding Rules (BER), an ITU and ISO standard developed by the ASN.1 consortium (ASN.1: ITU-T X.680, ISO/IEC 8824-1; BER: ITU-T X.690, ISO/IEC 8825-1).

ASN.1 and BER are used by widely adopted technologies and standards like:

- LDAP, Active Directory
- PKCS (Public Key Cryptography Standard)
- X.400 Electronic Mail
- Voice Over IP
- SNMP
- UMTS.

It is recommended to read the documents listed in the section “References” in order to be able to fully comprehend this documentation.

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY" and "OPTIONAL" in this document are to be interpreted as described in [\[RFC2119\]](#).

References

This document refers to the following specifications:

- BER (ITU-T X.690 08/2015) <https://www.itu.int/rec/T-REC-X.690-201508-I/en>
- ASN.1 (ITU-T X.680 07/2002) <http://www.itu.int/ITU-T/studygroups/com17/languages/X.680-0207.pdf>

Additional resources:

- Brief introduction to ASN.1 and BER from the Microsoft Knowledge Base <http://support.microsoft.com/kb/252648>
- "ASN.1 Complete", a very comprehensive and freely downloadable book covering ASN.1 and BER by Professor John Larmouth <http://www.oss.com/asn1/larmouth.html>

Basics

Overview

Ember encoded data must be decodable by any fully featured BER decoder. An Ember decoder on the other hand is not required to support the entire BER standard, since it is only required to implement the BER subset defined in this document. Therefore, an Ember encoder shall apply all the restrictions to the BER encoding defined in this document.

The Paragraph "Length Forms" and following define the most important parameters for the BER encoding that turn Ember into a mere subset of BER.

Tagging [X.690 8.14]

An encoded value always contains two tags, the outer or application tag, which defines the meaning or semantics of the data, and the inner tag which indicates the type of the encoded object. For types with primitive encoding (Boolean, Integer, Octet String, Real, UTF8String, Relative Object Identifier), the tagging environment shall be explicit without any exceptions. For types with constructed encoding (Sequence and Set), implicit tagging is often used for the inner tag. Besides the tag, the length of the encoded data is stored. So every object being encoded contains an outer Tag, an outer Length, an inner Tag, an inner Length and a Value, in short TLTLV. An encoding example can be found in the chapter "[The EmBER library](#)".

Length forms [X.690 8.1.3]

1. Generally, both length forms (definite and indefinite) are allowed; usage of the indefinite form though shall be permitted for containers only [X.690 8.1.3.2].
2. The length forms used by the inner and outer tag generated through explicit tagging must be identical. So if the outer tag specifies an indefinite length, the inner tag also has to specify an indefinite length.

Container Usage

The only containers that may be used are “Set” [X.690 8.11] and “Sequence” [X.690 8.10].

Basic document structure

The document must start with a container, and there shall only be one container at the top-level (like in the XML standard).

Types

Overview

The following BER types are supported by Ember:

1. Boolean
2. Integer
3. Real
4. UTF8String
5. Octet String
6. Set
7. Sequence
8. Relative Object Identifier

Boolean [X.690 8.2]

The decoder shall interpret 0 (zero) as “false” and any non-zero number as “true”. It is recommended for the encoder to use FF16 to encode “true”.

Integer [X.690 8.3]

Integers shall have a maximum size of 64 bits; the decoder may choose the data type to hold the decoded integer depending on the actual length. (E.g. 32 bit integer if the unsigned value is $< 2^{32}$, 64 bit otherwise).

Special care must be taken by the encoder to comply to rule [X.690 8.3.2]: the first 9 bits of the encoded integer value must not be the same.

To verify the correctness of the implementation, the following table should be used:

Input Value	Encoded result	
<i>Integers (base 10)</i>	<i>Length octets</i>	<i>Value octets (base 16)</i>
1	01	01
-1	01	FF

255	02	00 FF
127	01	7F
128	02	00 80
-128	01	80
$2^{16} - 1$	03	00 FF FF
2^{15}	03	00 80 00
-2^{15}	02	80 00

Real [X.690 8.5]

Real values must be encoded in binary form [X.690 8.5.6], base 2 [X.690 8.5.6.2]. Maximum precision shall be double (64 bit).

UTF8String [X.690 8.21]

The UTF8 encoding is supported.

Octet String [X.690 8.7]

This type shall be used to encode junks of binary data. The encoding must be primitive; splitting the octet string [X.690 8.7.3] is not permitted.

Relative Object Identifier [X.690 8.20]

A relative object identifier is a sequence of integers, which identify an object within a tree. The object identifier is encoded by writing the encoded length first, followed by all components encoded as multi-byte integers.

Set [X.690 8.11]

Every item in a “Set” container shall be implicitly optional; therefore an empty set is allowed. The tags of values in a Set container must be unique within the container context [X.680 26.3]. The order of values in a set shall be arbitrary [X.680 26.6].

Sequence [X.690 8.9, X.690 8.10]

This container shall represent the concept of either an ordered collection or a list type (Sequence-of [X.690 8.10]). The order of child values may be semantically significant [X.690 8.10.3]; therefore, tags do not necessarily have to be unique within the container context, though it is recommended to ensure tag uniqueness.

Implementation recommendations

Encoder

Ember puts a complex job upon the encoding entity: the encoder has to calculate the number of octets the encoding process will yield before the actual encoding is done. This is because the length octets are transmitted prior to the data octets. To make things worse, the tag and length octets preceding the data also vary in number. This makes it almost impossible to encode entire containers on the fly.

For encoding entities equipped with powerful hardware resources (like workstation and server PCs), it is therefore recommended to build the encoding in a DOM – like object tree, thereby calculating the length of each container node. This tree is easy to modify and quick to write out, but can possibly consume lots of memory.

For encoding entities that are based on less powerful hardware like embedded devices, hope lies in the generosity of the Basic Encoding Rules, which Ember does not confine:

- To encode containers, the encoder is free to use the indefinite length form instead of pre-calculating the number of value octets (which in turn contain each component's tag, length and value octets).
- The encoder is free to use the definite long length form [X.690 8.1.3.5] even for lengths smaller than 128. This allows fixing the number of octets used for length fields to – as an example – four.
- Depending on the DTD, a fixed number of octets can also be used to encode tags. E.g. if the DTD defines a maximum tag number of 127 (remember that tags are context-specific by default!), an encoder can always use two bytes to encode a tag – with the second byte containing the actual tag number and the five least significant bits of the first byte all set to “1”.
- Most primitive types inherently have a maximum length that is dependent on the encoder's hardware. E.g. for a 16 bit embedded system, the maximum number of octets an integer can be encoded into is usually 4 (for 32 bit “long” values). Thus, the value can be encoded into a fixed-length memory block, returning the number of bytes written.

Decoder

Decoding an Ember stream is quite straight-forward and the implementation should be the same for systems with powerful hardware resources as for embedded systems.

A decoder should implement one function to read and store the next item (tag, length and value) from the input stream. Additional functions should be implemented to convert the current value to the type indicated by the inner tag.

Stepping through the containers can either be implemented using recursion or by stacking container objects.

Glow specification

Introduction

This chapter describes the Glow specification, which is written in ASN.1 and called the Glow DTD (Design Type Document). The DTD defines several object types that are required for a device to represent its structure. The number of types is kept at a minimum so that it can be used for almost any kind of device. The most frequently used types are node and parameter. A node represents a device and its elements, like modules, any kind of extension cards or attached panels. Additionally, it may be used to introduce categories, like “Sources” and “Targets”.

A parameter then contains the properties of a control, like the range, the access (read, write, read-write, none), the name and so on.

Data types and properties

The tag format

As already mentioned, a tag may indicate the meaning and the type of the encoded data that follows next. A tag consists of two components, a class - which uses one byte - and a number, using up to four bytes. The combination of both defines the type.

The tag class accepts the following values:

Name	Value	Description
Universal	0x00	Predefined types. Should never be used.
Application	0x40	Application specific tags have the same meaning wherever seen and used.
Context	0x80	The meaning context specific tags depends on the location where they are seen
Private	0xC0	A special version of the context specific tag

If the tag defines a container type, like set or sequence, the 6th bit of the tag is set. Alternatively a binary OR operation can be performed (Class | 0x20). The meaning of the tag number depends on the class. If it is set to “Universal” it defines one of the default data types. The following table shows a list of standard BER types:

Number	Type		
1	Boolean	18	Numeric String
2	Integer	19	Printable String
3	Bitstring	20	Teletex String
4	Octet String	21	Videotex String
5	Null	22	IA5 String
6	Object Identifier	23	UTC Time
7	Object Descriptor	24	Generalized Time
8	External	25	Graphic String
9	Real	26	Visible String
10	Enumerated	27	General String
11	Embedded Pdv	28	Universal String
12	UTF8 String	29	Unspecified String
13	Relative Object	30	BMP String
16	Sequence	31	Last Universal
17	Set	0x80000000	Indicates an application defined type

Note that only the required subset of the universal types is supported in this library.

For example, a tag with the class set to “Universal” and the number set to 1 indicates that the encoded data represents a Boolean value.

Application defined tags

All other tag classes (Application and Context) may be used to define application specific types or, if used as (outer) application-tag, to indicate the usage. Each encoded object contains two tags, the first one identifying the meaning or semantics of the data or object, while the second one specifies the data type.

The Glow specification uses a set of attributes which are used to describe a node or parameter. Two common attributes are the identifier and the description string. The following section lists all attributes used by the Glow DTD. Since these attributes differ for each object type they are context specific. Please note that “Value” simply means that a property may be an Integer, Real,

UTF8String, an octet string value or a Boolean. The details about the “Value” type can be found in the next chapter.

Application defined types

In many cases the universal types suffice the requirements. Nevertheless it is possible to declare application defined types. Ember+ defines several types like Node, Parameter and Command.

The tag class of user defined types is set to “Application”-specific, while the number determines the type itself. The EmBER library provides a Type class which can be used to distinguish between universal and application defined types. The following section lists the types and their meaning used and defined by the Glow specification.

The complete ASN.1 notation of the Glow object types is listed in the chapter “[Ember+ ASN.1 Notation](#)”. Additionally, the chapter contains detailed descriptions about all objects for clarity.

Element

The element is a choice that contains one of the following types:

- Parameter
- Node
- Command
- StreamCollection
- Matrix (Glow 2.10 or higher)
- Function (Glow 2.20 or higher)

RootElement

- Element
- QualifiedNode
- QualifiedParameter
- QualifiedMatrix (Glow 2.10 or higher)
- QualifiedFunction (Glow 2.20 or higher)

The type an element contains can be determined by the type tag. A QualifiedNode and QualifiedParameter may only appear at root level, to reduce the complexity of the tree structure.

Value

A value is a choice that contains one of the following types:

- Integer
- Real
- UTF8String
- Boolean
- Octet String

Parameter

Type Tag: Application – 1

Type: Sequence

Members:

- Number
- Contents, which may contain:
 - Identifier
 - Description
 - Value
 - Minimum
 - Maximum
 - Access
 - Format
 - Enumeration
 - Factor
 - Online
 - Formula
 - Default
 - Command
 - StreamIdentifier
 - EnumMap
 - StreamDescriptor
- Children (must contain at most one Command)

A parameter represents an entity that at least contains a value.

When a data provider reports a parameter, the message must always contain the number. All other properties are optional and are listed within a set with Tag Context – 1. In general, automatic change notifications should only contain the properties that have changed, except when the parameter is marked as a stream. More behaviour rules can be found in the chapter “[Behaviour rules and other guidelines](#)”.

QualifiedParameter

Type Tag: Application – 9

Type: Sequence

Members:

- Path
- Contents
- Children (must contain at most one Command)

The contents are equal to the Parameter type. In comparison to the default Parameter object, the QualifiedParameter reports its complete path instead of its number. That way a parameter that is located in a large sub tree can be reported without the need to construct the whole tree structure, which saves CPU time and memory.

Command

Type Tag: Application – 2
Type: Sequence
Members: Number

A command may be appended to a node or a parameter as child element inside the ElementCollection property. A command only contains an integer identifying the command type, which may be [GetDirectory](#), [Subscribe](#) or [Unsubscribe](#).

Note: When the command is GetDirectory, an additional property called *DirFieldMask* can be provided. For details please see the [GetDirectory](#) command.

Node

Type Tag: Application – 3
Type: Sequence
Members:

- Number
- Contents, which may contain:
 - Identifier
 - Description
 - IsRoot
- Children

A node represents a device or one of its components. Like the parameter, it must contain a number which identifies the node while the session is active. All other properties are optional and therefore listed in a separate set.

QualifiedNode

Type Tag: Application – 10
Type: Sequence
Members:

- Path
- Contents
- Children

The contents are equal to the Node type. Compared to the default Node, the QualifiedNode reports its complete path, instead of only providing its number. That way, a node that is located in a big sub tree may also be reported at root level.

ElementCollection

Type Tag: Application – 4
Type: Sequence
Members: Entries of type Element

The ElementCollection is a sequence of Element types. A node or a parameter may contain an ElementCollection with commands or child nodes.

RootElementCollection

Type Tag: Application – 11
Type: Sequence
Members: RootElement

Root

App. Tag: Application – 0
Type: Choice of RootElementCollection or StreamCollection

StreamEntry

Type Tag: Application – 6
Type: Sequence
Members: StreamIdentifier, Value

Stream entries are used to report audio level data. Since these values change frequently, they may be transmitted via a different transport layer, like UDP. A StreamEntry is a sequence consisting of the unique stream-identifier and the current value. The StreamEntry contents must be transmitted in the specified order.

StreamCollection

Type Tag: Application – 5
Type: Sequence
Members: Entries of type StreamEntry

A StreamCollection is sequence of StreamEntry elements.

StringIntegerPair

Type Tag: Application – 7
Type: Sequence
Members: Name, Value

A tuple containing a enumeration name and an integer value. This type is used by the `StringIntegerCollection`, which is used by the `EnumMap` property.

StringIntegerCollection

Type Tag: Application – 8

Type: Sequence

Members: Entries of type [StringIntegerPair](#)

The `StringIntegerCollection` is a sequence that contains `StringIntegerPairs`. This structure is used by the enumeration map.

Glow specific properties

This chapter describes all properties that are used within this protocol.

Application specific properties

Root

Tag: Application – 0

Type: RootElementCollection, StreamCollection (Container)

Each ember message has to start with a RootElementCollection or the StreamCollection.

Default Element

Tag: Context – 0

Type: Usually used for a collection of nodes or parameters, where the tag is not used.

Node properties

All Node properties are context specific.

Number

Tag: Context – 0

Type: Integer

Each node must have a number that is unique within the current scope. It must not change while a session is active because it may be used to identify a node or parameter by using this number instead of its string identifier. The number must be equal to or greater than zero.

Contents

Tag: Context – 1

Type: Set

This set contains the optional properties of a node.

Children

Tag: Context – 2

Type: Sequence

The Children property is a sequence that contains the child-nodes and parameters of this node.

QualifiedNode properties

All QualifiedNode properties are context specific

Path

Tag: Context – 0

Type: Relative Object Identifier

The Path property contains the numeric path of this node.

Contents

Tag: Context – 1

Type: Set

This set contains the optional properties of a node.

Children

Tag: Context – 2

Type: ElementCollection

The Children property is a sequence that contains the child nodes and parameters of this node.

Node Contents**Identifier**

Tag: Context – 0

Type: UTF8String

This property contains the string identifier of a node. This name must be unique within the current scope, which means that all child nodes of one parent must have different names. The identifier of an entity must not change! Please also read the naming rules described in [“The identifier property”](#).

Description

Tag: Context – 1

Type: UTF8String

The Description contains the display name of a node. This property shall be displayed by a user interface if it is provided. Otherwise, the identifier should be used.

IsRoot

Tag: Context – 2

Type: Boolean

The IsRoot property indicates whether the current node is a root node or not. This flag may be used by providers acting as a kind of proxy, where several sub trees of different devices are merged into one tree.

IsOnline

Tag: Context – 3

Type: Boolean

Indicates whether the node is online or not. The default value is true. If a node is marked offline, all of its children must be considered offline as well.

SchemaIdentifiers

Tag: Context – 4

Type: Boolean

A single string comprised of schema identifiers separated by line feeds ('\n', 0). Each line identifies a schema that the sub-tree under this node complies with. See chapter [Ember+ 1.3: Schema Extensions](#).

Parameter properties

All Parameter properties are context specific.

Number

Tag: Context – 0

Type: Integer

Each parameter must have a number that is unique within the current scope. It must not change while a session is active because it may be used to identify a parameter by using this number instead of its string identifier. The number must be equal to or greater than zero.

Contents

Tag: Context – 1

Type: Set

This set contains the optional properties of a parameter.

Children

Tag: Context – 2

Type: ElementCollection

This collection may contain a command for the current parameter.

QualifiedParameter properties

All Parameter properties are context specific.

Path

Tag: Context – 0

Type: Relative Object Identifier

This property contains the numeric path of a parameter.

Contents

Tag: Context – 1

Type: Set

This set contains the optional properties of a parameter.

Children

Tag: Context – 2

Type: ElementCollection

This collection may contain a command for the current parameter.

Parameter Contents**Identifier**

Tag: Context – 0

Type: UTF8String

This property contains the string identifier of a parameter. This name must be unique within the current scope, which means that all child nodes and parameters of one parent must have different names. The identifier of an entity must not change!

Description

Tag: Context – 1

Type: UTF8String

Display name of a parameter. This property shall be displayed by a user interface if it is provided. Otherwise, the identifier should be used.

Value

Tag: Context – 2

Type: Value

This field stores the current value of a parameter. Supported types are Integer, Real and String. If the value contains an enumeration the type must be integer and the value contains the index of the enumeration entry to display, starting at index 0.

Sending a value implies that the node is online. A consumer which receives a value should mark this node and its parents as online. A provider **MUST NOT** send values in a tree node which is actively held offline by itself or any of its parent.

Minimum

Tag: Context – 3

Type: Value

The Minimum property contains the smallest value allowed, which may either be a real or integer value.

Maximum

Tag: Context – 4

Type: Value

The Maximum property contains the largest value allowed, which may either be a real or integer value.

Access

Tag: Context – 5

Type: Integer

Indicates how this parameter can be accessed. The following values are available:

- None (0)
- Read Only (1)
- Write Only (2)
- Read/Write (3)

Format

Tag: Context – 6

Type: UTF8String

This is an optional format string. This property may contain a C-Style format string which may be used to append a unit or define the number of digits that should be displayed. To simplify implementation of creating a format string on consumer side, only the format-characters d, f and x must be used. For example, "%.2f°\ndB" or "%d Units" are valid format strings. The ° character should be used to indicate where the unit string begins.

Enumeration

Tag: Context – 7

Type: UTF8String

A single string containing the values of an enumeration, separated by a line feed ('\n', 10). The values must internally be enumerated from 0 to N. If an entry shall not be displayed in a user interface, the entry must begin with '~'.

Factor

Tag: Context – 8

Type: Integer

This property may be used if a device is not able to process decimal values. It may then provide a factor instead. The consumer then has to divide the reported value when it displays it and multiply it with this factor when it wants to change the parameter. The type of the factor is integer.

IsOnline

Tag: Context – 9

Type: Boolean

This field contains the online state. When a node reports an offline state, the consumer should assume that all child nodes are offline as well. On the other side, when a node reports an online state, only the parents of this node may be changed to online.

Formula

Tag: Context – 10

Type: UTF8String

The Formula is an optional mathematical term. The device must provide two formulas separated by a linefeed character ('\n', 10). The first formula is used to transform the device value into a display value (provider to consumer), while the second one is used transform it back to a device value again (consumer to provider).

Note

Please note that there are two formats supported:

- A term can be provided as RPN term, without a character that represents a variable for the current device value. An example would be 5*, which would multiply the current device value with 5. The device value always has to be pushed on the computation stack first.
- When the expression is enclosed in parentheses a term is interpreted as mathematical expression, including the \$ character as placeholder for the current device value. A valid term would be (5 * \$).
A detailed description of all supported functions can be found in the document “Ember+ Formulas”.

The RPN variant is provided for backward compatibility.

Step

Tag: Context – 11

Type: Integer

This property specifies the step to use when incrementing or decrementing a value. Using this property should be avoided, since it causes problems when using real values. Instead the **Factor** should be used in combination with a value of type integer.

This property may be removed in future releases of the Ember+ libraries.

Default

Tag: Context – 12

Type: Value

Default value of a parameter.

Type

Tag: Context – 13

Type: Integer

The Type can be used to provide a hint about the value type of the parameter. The following values are currently available:

- Integer (1)
- Real (2)
- String (3)
- Boolean (4)
- Trigger (5)
- Enum (6)
- Octets (7)

StreamIdentifier

Tag: Context – 14

Type: Integer

A number used to identify an audio level meter. Since peak meter data is usually reported frequently (like every 80ms) it may be necessary to use a separate transport layer like UDP in order to transmit their values. That's why audio streams require a globally unique identifier which must not be negative. Their values are being reported in a separate container, the StreamCollection.

EnumMap

Tag: Context – 15

Type: StringIntegerCollection

This field provides a more complex description of an enumeration. The StringIntegerCollection is a sequence of a StringIntegerPair which is tuple name, value tuple. This field can be used if a device is unable to use a consecutive enumeration.

StreamDescriptor

Tag: Context – 16

Type: StreamDescription

A StreamDescriptor is required when a single stream entry contains more than one value and is provided as octet string. The descriptor then provides the internal data offset and the stream value format. The formats available can be found in the [DTD](#) at the end of the document.

SchemalIdentifiers

Tag: Context – 17

Type: Boolean

A single string comprised of schema identifiers separated by line feeds ('\n', 0). Each line identifies a schema that the parameter complies with. See chapter [Ember+ 1.3: Schema Extensions](#).

Application defined commands

In addition to these properties there is also a small list of commands. Commands are only used by the consumer, the provider always only has to respond or report its current status. The command type is determined by the number of a command, which must always be provided.

GetDirectory

Number: 32

This command may be executed on a node and a parameter. It requests all child nodes and parameters of the node containing this command. The response shall also include all attributes of the reported entities.

The purpose of this command is to obtain the complete or only a part of the structure of a data provider.

Additionally, this command implicitly indicates that a consumer is interested in the entities below the node the command has been sent for. That means, that parameters belonging to this node should report value changes to the consumer automatically.

When a consumer performs a GetDirectory request on a node without any children, the node must report itself without any properties set (even the identifier has to be omitted). Otherwise it is not possible to distinguish it from an ordinary update. That way, the consumer knows that a query has been performed on this node and that it has no child elements.

This command can also be applied to a parameter to query its properties.

Besides, a *dirFieldMask* can also be provided, which is tagged with “Context – 1”. This mask determines the properties a consumer is interested in:

- All (-1)
- Default (0)
- Identifier (1)
- Description (2)
- Value (4)

All and *Default* should return the same result, which is a node or parameter with all its properties set. When a consumer requests a property that doesn't exist, the provider should report the node or parameter without any properties.

Subscribe

Number: 30

This command is used to subscribe to a parameter which has a [StreamIdentifier](#). Parameters without `StreamIdentifier` should report their value changes automatically (at the latest when a parameter has been queried by `GetDirectory` command), while a parameter that has a `StreamIdentifier` only transmits its changes when it is subscribed. The main purpose of the subscription mechanism is to reduce network traffic and the CPU load for devices with limited resources. A common example is the transmission of audio levels for multiple channels. These values usually change frequently (several times per second) and in many cases a consumer is not interested in receiving these updates all the time.

So for parameters that change often, a `StreamIdentifier` should be used to avoid that the provider has to encode and transmit the updated values all the time, although no one is currently interested in receiving them. When a `StreamIdentifier` is set for a parameter, a consumer must subscribe to this parameter. That way, the provider knows that someone is interested in receiving value updates and transmits them. But these updates are not being sent in the usual way, these values are transported within a [StreamCollection](#) container instead, which contains a list of [StreamEntries](#).

Unsubscribe

Number: 31

This command is used to unsubscribe from a node or parameter. When applied to a node, it indicates that the consumer is no longer interested in receiving value updates from the parameters that are located below this node. In that case, all parameter subscription for that consumer must be removed.

When applied to a single parameter, only the existing subscription for the consumer that sent the request must be removed.

Ember+ 1.1: Matrix Extensions

Introduction

This chapter describes how Ember+ 1.1 integrates matrices into the Glow schema, which for Ember+ 1.1 has the version 2.10. A matrix in the sense of this chapter is a two-dimensional array of Boolean values.

Possible applications of matrices in the context of gadget control:

- Signal routing
- GP-I/O signaling
- Key assignment for Intercom systems
- Group and conference management for Intercom systems
- Bus assignments of summation matrices (e.g. Mixing Consoles)

The entities represented by the rows and columns of a matrix are called signals. The signals represented by the matrix columns are called targets. The signals represented by the matrix rows are called sources.

The values contained in the matrix cells may be interpreted as connections from sources to targets: If `matrix[S, T]` is true, source S is connected to target T.

Matrices inherently have three different connection semantics:

- 1:N
one source may be connected to n targets, but each target must not be connected to more than one source
- 1:1
one source may be connected to only one target, and each target must not be connected to more than one source
- N:N
a source may be connected to n targets, and a target may have n sources connected to it.

Ember+ must support all three matrix types.

For N:N matrices, connections are usually described by multiple values – not only the Boolean state of the connection, but also a ratio or weight for each connected source. This proposal takes this into account by allowing a matrix to refer to a collection of Parameters describing each signal and connection.

Also, the signals making up a matrix may have multiple labels, each of which may be editable. This proposal takes this into account by allowing a matrix to refer to multiple collections of Parameters which describe signal labels.

Type Definitions

Glow 2.10 introduces several new application-defined types, the most important of which is called Matrix.

This new type stands alongside the types `Node` and `Parameter` (which are already defined in Glow 2.5).

Matrix

- `number [0] Integer32`
The number of the matrix object. Has same semantics as `Parameter.number` and `Node.number`.
- `contents [1] MatrixContents OPTIONAL`
Contents set of the matrix object. Analogical to `Parameter.contents` and `Node.contents`.
- `children [2] ElementCollection OPTIONAL`
Contains child elements of the matrix object. Conceptually, matrix objects are considered leaves in the Glow tree. Though, a matrix may contain a Glow Command in the same way as a parameter.
- `targets [3] TargetCollection OPTIONAL`
A collection of `Target` objects that fill columns of the matrix. This property must only be present if the matrix addressing mode is non-linear. A linear matrix may use this property to restrict connections to the targets listed under this property.
- `sources [4] SourceCollection OPTIONAL`
A collection of `Source` objects that fill rows of the matrix. This property must only be present if the matrix addressing mode is non-linear. A linear matrix may use this property to restrict connections to the sources listed under this property.
- `connections [5] ConnectionCollection OPTIONAL`
A collection of `Connection` objects that describe each target's active connection. The provider must report one `Connection` object for each target.

When the children of a node are requested (using the `GetDirectory` command) and this node contains a `Matrix` object, the matrix is transmitted only containing the properties `number` and `contents`.

The consumer may then issue a `GetDirectory` command as a child of the matrix object. The provider must reply with the complete matrix object, containing at least the property `connections`. In the case of a matrix with non-linear addressing mode, the `targets` and `sources` properties are also required.

As soon as a consumer issues a `GetDirectory` command on a matrix object, it implicitly subscribes to matrix connection changes. The consumer may afterwards use the `Unsubscribe` command to cancel this subscription.

MatrixContents

- `identifier [0] EmberString`
The object identifier. Analogical to `Parameter.identifier` and `Node.identifier`.
- `description [1] EmberString OPTIONAL`
The display name of the matrix object. Analogical to `Parameter.description` and `Node.description`.

- **type** [2] **MatrixType** OPTIONAL
The type of the matrix: either 1:N, 1:1 or N:N. If this property is not present, 1:N is assumed.
- **addressingMode** [3] **MatrixAddressingMode** OPTIONAL
The addressing mode used by the matrix. Either linear (signal numbers are row/column indices) or non-linear (signal numbers are random).
- **targetCount** [4] **Integer32**
If **addressingMode** is linear, this defines the number of matrix columns. If **addressingMode** is non-linear, this defines the number of targets contained in the **targets** property.
- **sourceCount** [5] **Integer32**
If **addressingMode** is linear, this defines the number of matrix rows. If **addressingMode** is non-linear, this defines the number of sources contained in the **sources** property.
- **maximumTotalConnects** [6] **Integer32** OPTIONAL
This property is only significant for N:N matrices. It defines the maximum number of active connections for the entire matrix. If this property is not present, **targetCount** * **sourceCount** is assumed.
- **maximumConnectsPerTarget** [7] **Integer32** OPTIONAL
This property is only significant for N:N matrices. It defines the maximum number of sources that may be connected to one target. If this property is not present, **sourceCount** is assumed.
- **parametersLocation** [8] **ParametersLocation** OPTIONAL
This property may refer to the location of parameters associated with the signals and connections of the matrix.
- **gainParameterNumber** [9] **Integer32** OPTIONAL
This property is only significant for N:N matrices and may only be present if the **parametersLocation** property is also present. It may contain the number of the parameter that describes the gain or ratio of a connection. This parameter must exist for every connection under the location referred to by the **parametersLocation** property. It must be either of type **INTEGER** or of type **REAL**.
- **labels** [10] **LabelCollection** OPTIONAL
This property may contain multiple references to Nodes which contain signal label parameters. These parameters must be of type **EmberString**.
- **schemaIdentifiers** [11] **EmberString** OPTIONAL
A single string comprised of schema identifiers separated by line feeds ('\n', 0). Each line identifies a schema that the matrix complies with. See chapter [Ember+ 1.3: Schema Extensions](#).

MatrixType

- **oneToN** (0)
Default. The matrix uses 1:N connection semantics.
- **oneToOne** (1)

The matrix uses 1:1 connection semantics.

- nToN (2)

The matrix uses N:N connection semantics.

MatrixAddressingMode

- linear (0)

Default. The matrix uses linear addressing mode.

- nonLinear (1)

The matrix uses non-linear addressing mode.

ParametersLocation


- basePath RELATIVE-OID

Contains an absolute path to a node under which parameters associated with signals or connections reside.

- inline Integer32

Contains a single sub-identifier to be appended to the path of the matrix. The resulting path refers to the node under which parameters associated with signals or connections reside. This node is not reported in a response to a **GetDirectory** command, since huge matrices with hundreds or thousands of connections will not be able to publish all connection and signal parameters in the Glow tree. The provider must create the parameter objects on-the-fly when a **GetDirectory** command is issued to a sub-node containing the parameters for a specific signal or connection.

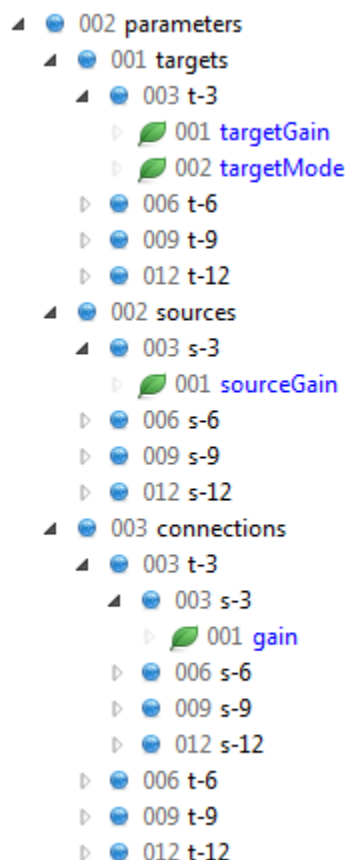
The sub-tree under the node a **ParametersLocation** object refers to must be structured like this:



In this sample, a **ParametersLocation** object refers to the node “2 - parameters”. This node must contain three well-known nodes: “1 - targets”, “2 - sources” and “3 - connections”. The node “1 - targets” must contain one node for each target associated with parameters. Each of these nodes must have the number of the corresponding target.

The node “2 – sources” must contain one node for each source associated with parameters. Each of these nodes must have the number of the corresponding source.

The node “3 – connections” must contain one node for each target. These target nodes must themselves contain one node for each source:



To request all parameters associated with target 3, the consumer may issue a [GetDirectory](#) command contained in a [QualifiedNode](#) object with the path “<parametersLocation>.1.3”.


To request all parameters associated with source 6, the consumer may issue a [GetDirectory](#) command contained in a [QualifiedNode](#) object with the path “<parametersLocation>.2.6”.

To request all parameters associated with connection “target 9 <- source 12”, the consumer may issue a [GetDirectory](#) command contained in a [QualifiedNode](#) object with the path “<parametersLocation>.3.9.12”.

Label

- `basePath [0] RELATIVE-OID`
This property contains an absolute path to a node under which label parameters associated with signals reside.
- `description [1] EmberString`
This property contains a free-text description of the label, which may be used by consumers to let the user choose the label to display.

The sub-tree under the node a [Label1.basePath](#) refers to must be structured like this:



In this sample, a `Label.basePath` refers to the node “1 - labels”. This node must contain two well-known nodes: “1 - targets” and “2 - sources”.

The node “1 - targets” must contain one parameter of type `EmberString` for each target associated with a label. Each of these parameters must have the number of the corresponding target.

The node “2 - sources” must contain one parameter of type `EmberString` for each source associated with a label. Each of these parameters must have the number of the corresponding source.

Signal

- number [0] Integer32

This property contains the number of the signal.

If the addressing mode of the encompassing matrix is linear, this number is the zero-based index of the row or column containing the signal.

If the addressing mode of the encompassing matrix is non-linear, this is a random number uniquely identifying the signal.

Connection

- target [0] Integer32

This property contains the number of the connection target.

- sources [1] PackedNumbers OPTIONAL

This property contains the numbers of all sources connected to `target`. If not present, the target does not have any active connections.

- operation [2] ConnectionOperation OPTIONAL

This property may contain a hint on how the set of sources must be interpreted. If this property is not present, “absolute” is assumed.

- disposition [3] ConnectionDisposition OPTIONAL

When the connection object is sent by the provider as the response to a connect request, this property may contain an error code indicating the execution state of the operation. If this property is not present, “tally” is assumed.

The consumer should not transmit this property at all when issuing connect requests.

PackedNumbers

This type is an alias for the primitive ASN.1 type RELATIVE-OID. It is used to hold a vector of integer values in a compact way.

ConnectionOperation

- absolute (0)
Default. This value indicates that the list of source numbers in the sources property of the encompassing Connection object is absolute. When a provider tallies a connection, it must always use this value, independent of the value of the disposition property.
- connect (1)
This indicates that the sources in the sources property of the encompassing Connection object should be connected to the given target. Only used by the consumer when issuing a connect request.
- disconnect (2)
This indicates that the sources in the sources property of the encompassing Connection object should be disconnected from the given target. Only used by the consumer when issuing a connect request.

ConnectionDisposition

- tally (0)
Default. The operation property of the encompassing Connection object must be set to “absolute” and the sources property must contain the absolute set of sources currently connected. This error level is usually only used when the list of connections is transmitted in the response to [GetDirectory](#) command.
- modified (1)
This error level may be used by the provider to indicate that the connection has changed. The sources property of the encompassing Connection must contain the absolute set of sources currently connected.
- pending (2)
May be used by the provider to signal that the connect operation was queued, but is not yet current.
- locked (3)
May be used by the provider to signal that the connect operation could not be executed because the target is locked. The sources property of the encompassing Connection must contain the absolute set of sources currently connected.

QualifiedMatrix

- path [0] RELATIVE-OID
The absolute path to the Matrix. Uses the same semantics as [QualifiedParameter.path](#) and [QualifiedNode.path](#).
- contents [1] MatrixContents OPTIONAL
See [Matrix.contents](#).
- children [2] ElementCollection OPTIONAL
See [Matrix.children](#).

- targets [3] TargetCollection OPTIONAL
See [Matrix.targets](#).
- sources [4] SourceCollection OPTIONAL
See [Matrix.sources](#).
- connections [5] ConnectionCollection OPTIONAL
See [Matrix.connections](#).

Use Cases

The following use cases are illustrated as XML samples (Glow is easily transposed to XML).

Transmission of the Matrix Description

Consumer -> Provider:

```
<QualifiedNode path="1.2">
  <children>
    <Command number="32" />
  </children>
</QualifiedNode>
```

This message requests all children under the node with path "1.2".

Provider -> Consumer:

```
<QualifiedNode number="1.2">
  <children>
    <Matrix number="1">
      <contents>
        <identifier>matrix</identifier>
        <description>Sample Matrix</description>
        <targetCount>4</targetCount>
        <sourceCount>4</sourceCount>
        <type>nToN</type>
        <parametersLocation type="RelativeOid">1.2.2</parametersLocation>
        <gainParameterNumber>1</gainParameterNumber >
        <labels>
          <Label basePath="1.2.3.1" description="Primary" />
          <Label basePath="1.2.3.2" description="Internal" />
        </labels>
      </contents>
    </Matrix>
  </children>
</QualifiedNode>
```

This message reports one child under the requested node: a matrix of type N:N, with linear addressing mode, 4 targets, 4 sources, a reference to associated parameters and references to two layers of labels.

Transmission of Connections

Consumer -> Provider:

```
<QualifiedMatrix path="1.2.1">
  <children>
    <Command number="32" />
  </children>
</QualifiedMatrix>
```

This message requests all sub-elements of the matrix (targets, sources and connections).

Provider -> Consumer:

```
<QualifiedMatrix path="1.2.1">
  <contents>
    <identifier>matrix</identifier>
    <!-- ... -->
  </contents>
  <connections>
    <Connection target="0">
      <sources>3</sources>
    </Connection>
    <Connection target="1">
      <sources>0.1</sources>
    </Connection>
    <Connection target="2">
      <sources>3.1.2</sources>
    </Connection>
    <Connection target="3" />
  </connections>
</QualifiedMatrix>
```

This message includes the matrix contents and connections. Targets and sources are omitted since the matrix uses linear addressing mode.

Issuing Connects

Consumer -> Provider:

```
<QualifiedMatrix path="1.2.1">
  <connections>
    <Connection target="0">
      <sources>0.2</sources>
      <operation>connect</operation>
    </Connection>
    <Connection target="1">
      <sources>2</sources>
      <operation>connect</operation>
    </Connection>
  </connections>
</QualifiedMatrix>
```

This message requests a connection to be made from sources 0 and 2 to target 0 and a connection to be made from source 2 to target 1.

Provider -> Consumer:

```
<QualifiedMatrix path="1.2.1">
  <connections>
    <Connection target="0">
      <sources>3.0.2</sources>
      <disposition>taken</disposition >
    </Connection>
    <Connection target="1">
      <sources>0.2</sources>
      <disposition>taken</disposition>
    </Connection>
  </connections>
</QualifiedMatrix>
```

This message reports the new state of the connections to targets 0 and target 1.

Issuing Disconnects

Consumer -> Provider:

```
<QualifiedMatrix path="1.2.1">
  <connections>
    <Connection target="0">
      <sources>3</sources>
      <operation>disconnect</operation>
    </Connection>
  </connections>
</QualifiedMatrix>
```

This message requests source 3 to be disconnected from target 0.

Provider -> Consumer:

```
<QualifiedMatrix path="1.2.1">
  <connections>
    <Connection target="0">
      <sources>0.2</sources>
      <disposition>taken</disposition>
    </Connection>
  </connections>
</QualifiedMatrix>
```

This message reports the new state of the connections to target 0.

Requesting Labels

Consumer -> Provider:

```
<QualifiedNode path="1.2.3.1.1">
```



```

<!-- path: 1.2.3.1      .1
      label.basePath targets -->
<children>
  <Command number="32" dirFieldMask="4" />
</children>
</QualifiedNode>

```

This is a usual **GetDirectory** message issued on the first layer of target labels reported by the matrix. It also asks the provider to only return the values of the found parameters.

Provider -> Consumer:

```

<QualifiedNode path="1.2.3.1.1">
  <children>
    <Parameter number="0">
      <contents>
        <value type="UTF8">Primary Label of Target 0</value>
      </contents>
    </Parameter>
    <Parameter number="1">
      <contents>
        <value type="UTF8">Primary Label of Target 1</value>
      </contents>
    </Parameter>
    <Parameter number="2">
      <contents>
        <value type="UTF8">Primary Label of Target 2</value>
      </contents>
    </Parameter>
    <Parameter number="3">
      <contents>
        <value type="UTF8">Primary Label of Target 3</value>
      </contents>
    </Parameter>
  </children>
</QualifiedNode>

```

This response includes all label parameters. Note that each parameter's number equals the number of the target it is associated with.

Changing Labels

Consumer -> Provider:

```

<QualifiedParameter path="1.2.3.1.1.3">
  <!-- path: 1.2.3.1      .1      .3
      label.basePath targets target3 -->
  <contents>
    <value type="UTF8">New Primary Label of Target 3</value>
  </contents>
</QualifiedNode>

```

This is the standard procedure for changing a parameter.

Provider -> Consumer:

```
<QualifiedParameter path="1.2.3.1.1.3">
  <contents>
    <value type="UTF8">New Primary Label of Target 3</value>
  </contents>
</QualifiedNode>
```

And this is the response signaling success.

Requesting Connection Parameters

Consumer -> Provider:

```
<QualifiedNode path="1.2.2.3.0.3">
  <!-- path: 1.2.2          .3          .0          .3
           parametersLocation connections target0 source3 -->
  <children>
    <Command number="32" />
  </children>
</QualifiedNode>
```

This is a usual `GetDirectory` message issued on the node including the parameters for connection `target0 <- source3`.

Provider -> Consumer:

```
<QualifiedNode path="1.2.2.3.0.3">
  <children>
    <Parameter number="1">
      <contents>
        <identifier>gain</identifier>
        <description>gain of connection target0 <- source3</description>
        <value type="Real">-64.0</value>
        <minimum type="Real">-128.0</value>
        <maximum type="Real">15.0</maximum>
        <access>readWrite</access>
      </contents>
    </Parameter>
    <Parameter number="2">
      <contents>
        <identifier>peak</identifier>
        <description>peak level of connection target0 <-
source3</description>
        <type>Integer</type>
        <minimum type="Real">0</value>
        <maximum type="Real">255</maximum>
        <streamIdentifier>110</streamIdentifier>
      </contents>
    </Parameter>
  </children>
</QualifiedNode>
```

This response includes two parameters: a gain parameter and a streaming-enabled peak parameter.

Changing Connection Parameters

Consumer -> Provider:

```
<QualifiedParameter path="1.2.2.3.0.3.1">
  <!-- path: 1.2.2      .3      .0      .3      .1
           parametersLocation connections target1 source4 gain -->
  <contents>
    <value type="Real">10.0</value>
  </contents>
</QualifiedParameter>
```

This is the standard procedure for changing a parameter.

Provider -> Consumer:

```
<QualifiedParameter path="1.2.2.3.0.3.1">
  <contents>
    <value type="Real">10.0</value>
  </contents>
</QualifiedParameter>
```

And this is the response signaling success.

Performance Characteristics

The following section lists common use cases and the number of bytes needed for transmission. All byte lengths take into account framing and escaping.

- Response to **GetDirectory** command on 1:N matrix 200x200 (transmits matrix contents, 200 targets, 200 sources, 200 connections): **6761 Bytes**.
- Setting a single connection on same matrix (transmits 1 connection): **46 Bytes**.
- Reporting a single connection with **disposition "modified"** on same matrix (transmits 1 connection): **51 Bytes**.
- Response to **GetDirectory** command on N:N matrix 4x4 with 4 active connections (transmits matrix contents, 4 targets, 4 sources, 4 connections): **247 Bytes**.
- Response to **GetDirectory** command on same matrix but with 16 connections: **259 Bytes**.
- Response to **GetDirectory** command on N:N matrix 1000x1000 with 1000 active connections: **36517 Bytes**.
- Same response as above, but only connections (**GetDirectory** command refined with **dirFieldMask**): **16211 Bytes**.
- Response to **GetDirectory** command on same matrix but with hypothetical maximum of 1000000 (one million) active connections: **2025838 Bytes** (1.93 Mbytes).
- Reporting a connection of one target to 1000 sources on same matrix (transmits one connection): **2051 Bytes**.

Ember+ 1.2: Function Extensions

Introduction

Although the main purpose of Ember+ is describing and mutating the state of interconnected systems by encoding the state itself, there are use cases which require remote procedure calls through the Ember+ interface. In some cases, the old way of triggering actions on the remote host through usage of simple “trigger” parameters is not sufficient because of the lack of trigger arguments.

Therefore this section describes how function calls are integrated into the Glow schema, which for Ember+ 1.2 has the version 2.20.

Please note: functions should only be used to add transaction functionality to a tree that otherwise describes the provider’s state by using parameters. Also, a provider must be aware that a consumer (or a cluster of redundant consumers) may invoke a function multiple times even if only one invocation was triggered by the user. This implies that a function should never change the provider’s state relatively, but only absolutely – e.g. “`goToPosition(int absolutePosition)`” instead of “`changePositionBy(int relativeDelta)`”.

Functions shall be described as leafs in the Glow tree (like parameters). The description shall include the number, names and types of function arguments and the type definition of the function’s return value. Both argument list and return value shall be described as n-tuples of primitive types.

Functions shall be invoked through a new Command (“invoke” command): the consumer may send a Function object containing an invoke command, which in turn contains the argument values and an invocation id. The provider shall respond with an object of the new type `InvocationResult`, which repeats the invocation id issued by the consumer and contains the return value.

The consumer may issue “Fire and forget” function calls by omitting the invocation id. This shall only be allowed for functions which do not yield a result.

Type Definitions

Function

- `number` [0] `Integer32`,
The number of the Function object. Has same semantics as `Parameter.number` and `Node.number`.
- `contents` [1] `FunctionContents` OPTIONAL
Contents set of the function object. Analogical to `Parameter.contents` and `Node.contents`.
- `children` [2] `ElementCollection` OPTIONAL

Contains child elements of the function object. Conceptually, function objects are considered leaves in the Glow tree. Though, a function may contain a Glow Command in the same way as a parameter or matrix.

QualifiedFunction

- **path** [0] RELATIVE-OID
The absolute path to the Function. Uses the same semantics as QualifiedParameter.path and QualifiedNode.path.
- **contents** [1] FunctionContents OPTIONAL
see Function.contents
- **children** [2] ElementCollection OPTIONAL
see Function.children

FunctionContents

- **identifier** [0] EmberString OPTIONAL
The object identifier. Analogical to Parameter.identifier and Node.identifier.
- **description** [1] EmberString OPTIONAL
The display name of the Function object. Analogical to Parameter.description and Node.description.
- **arguments** [2] TupleDescription OPTIONAL
Describes the arguments the function takes with name and type. If not present, the function does not take any arguments.
- **result** [3] TupleDescription OPTIONAL
Describes the return value of the function, which is a tuple of arbitrary length and type.

TupleItemDescription

- **type** [0] P
The type of the argument. Only primitive types are supported.
- **name** [1] EmberString OPTIONAL
The name of the argument. May be omitted if the function only takes one argument.

Invocation

- **invocationId** [0] Integer32 OPTIONAL
An integer number generated by the consumer and repeated by the provider in the InvocationResult. Used by the consumer to relate InvocationResults to Invocations.
- **arguments** [1] Tuple OPTIONAL
The values of the arguments as described in Function.arguments.

InvocationResult

- **invocationId** [0] Integer32
The invocation id originally issued by the consumer in the invoke command.
- **success** [1] BOOLEAN OPTIONAL
True or omitted if the provider did not encounter any errors while executing the function.

- **result** [2] Tuple OPTIONAL
The return value of the function call. The values in this tuple must match the description given in `FunctionContents.result`.

Use Cases

The following use cases are illustrated as XML samples (Glow is easily transposed to XML).

Transmission of the Function Description

Consumer -> Provider:

```
<QualifiedNode path="1.2">
  <children>
    <Command number="32" />
  </children>
</QualifiedNode>
```

This message requests all children under the node with path "1.2".

Provider -> Consumer:

```
<QualifiedNode path="1.2">
  <children>
    <Function number="1">
      <contents>
        <identifier>setObjectNameRecursive</identifier>
        <description>Sample Function</description>
        <arguments>
          <TupleItemDescription type="INTEGER" name="objectId" />
          <TupleItemDescription type="BOOLEAN" name="isRecursive" />
          <TupleItemDescription type="UTF8" name="newName" />
        </arguments>
        <result>
          <TupleItemDescription type="integer" name="changeCount" />
        </result>
      </contents>
    </Function>
  </children>
</QualifiedNode>
```

This message reports one child under the requested node: a Function with three arguments (int objectId, bool isRecursive, string newName) and returning the number of changes committed by calling the function.

Invocation of Functions

Consumer -> Provider:

```
<QualifiedFunction path="1.2.1">
```

```
<children>
  <Command number="33">
    <invocation invocationId="1">
      <arguments>
        <Value type="INTEGER">123</Value>
        <Value type="BOOLEAN">True</Value>
        <Value type="UTF8">Herbert</Value>
      </arguments>
    </invocation>
  </Command>
</children>
</QualifiedFunction>
```

This message issues a recursive rename of all objects beneath object 123.

Provider -> Consumer:

```
<InvocationResult invocationId="1">
  <result>
    <Value type="INTEGER">74</Value>
  </result>
</InvocationResult>
```

This message reports that 74 objects have been changed through invocation with id 1.

Ember+ 1.3: Schema Extensions

Purpose

The purpose of schema definitions is to provide a way for automated systems to extract formal type information for sub9trees within an Ember+ tree, without having to rely on structural subtyping. With type information readily available, a consumer may provide high level functionality utilizing the identifiable parts of the tree on otherwise unknown Ember+ trees. In contrast to specialized extensions of the Ember+ DTD for certain domains, schemas have the advantage that they may be defined ad hoc by the implementer of an Ember+ provider without obstructing the utility of the data for consumers that have no specialized knowledge of the schema in question.

Identification within an Ember+ Tree

Schema identifiers may occur anywhere within an Ember+ tree as part of a node, as part of a parameter or as part of a matrix definition. A schema identifier is a string that refers to a schema using reverse domain name notation (e.g. `de.1-s-b.emberplus.schema1`).

Nodes, parameters and matrices may comply with multiple schemas, as long as these schemas do not define conflicting identifiers.

Formally the schema identifier must conform to the syntax rules specified in RFC1035, section 2.3.1, while explicitly disregarding the extensions specified in RFC2673 and RFC4343. In addition the following restriction is made with regard to the syntax rules specified in RFC1035:

```
<letter> ::= "a" | ... | "z"
```

Scope

When a schema identifier is encountered within an Ember+ tree its scope spans the entire sub9tree rooted at the element, as part of which the schema identifier is encountered, although a schema definition is free to restrict the scope as it sees fit. A schema definition is not allowed to extend the scope beyond the sub9tree to siblings or parents of the element, as part of which the schema identifier is encountered.

Schema definition

The means through which a schema is defined is up to the implementer. It should formally and unambiguously define the schema, by

1. Defining to which type of element the schema may be applied.
2. Constraining the set of direct or indirect children that may or must appear below the element to which the schema applies.

3. Formally describing the meaning of certain parameters within the application domain, to which the sub-tree relates
4. Referring to other schemas to which certain child elements must conform.

This list is of course non exhaustive and any means required to properly describe the schema are acceptable.

Even though it is valid for a schema definition to restate certain constraints that may also be expressed using the means provided by the Ember+ DTD, a schema must not be used to obsolete or omit such information, because even in the presence of schemas generic Ember+ consumers that have no further knowledge of a specific schema must not be hindered in their correct operation.

Examples of this are:

- Even if a schema definition specifies a valid min/max range for a certain parameter the minimum and maximum attributes must still be set accordingly.
- Even if a schema definition specifies how an internal value has to be formatted for presentation purposes the formula attribute must still be set accordingly.

Application Notes

- A node governed by a schema application is always allowed to contain more parameters or children than those specified in the schema definition, as long their presence is not explicitly forbidden.
- Commands do not allow the application of a schema, because commands are considered standard operations that are mandated and defined by the Ember+ protocol specification.
- Functions do not allow the application of a schema, because they are considered operations within the context of their parent node (like methods in an object oriented context) and are thus to be defined as part of a schema, whose application governs their parent element.
- If future extensions to a schema definition are likely, we suggest post-fixing the schema identifier with a version identifier such as .vXXXX.YYYY. Where XXXX denotes the major version number and YYYY denotes the minor version number, the latter of which may be omitted. Changes in the minor version number are expected to be backwards compatible, whereas a different major version number must be considered incompatible by consumers that are not aware of the number encountered.

Syntax of Schema Identifiers

This BNF is taken from RFC1035 with the modification mentioned above, and modified to not allow a single space character as an empty schema identifier.

```
<schema-ident> ::= <label> | <schema-ident> "." <label>
<label> ::= <letter> [ [ <ldh-str> ] <let-dig> ]
<ldh-str> ::= <let-dig-hyp> | <let-dig-hyp> <ldh-str>
```

```
<let-dig-hyp> ::= <let-dig> | "-"  
<let-dig> ::= <letter> | <digit>  
<letter> ::= "a" | ... | "z"  
<digit> ::= "0" | ... | "9"
```

References

- RFC1035: Domain names – implementation and specification
<http://tools.ietf.org/html/rfc1035>
- RFC2673: Binary Labels in the Domain Name System
<http://tools.ietf.org/html/rfc2673>
- RFC4343: Domain Name System (DNS) Case Insensitivity Clarification
<http://tools.ietf.org/html/rfc4343>

Ember+ 1.4: Templates

Purpose

Ember+ trees often consist of hierarchical sets of structurally equivalent elements or sub-trees. While the Ember+ Schema extension introduced in Glow DTD v2.30 allows easy identification of certain sub-trees through a form of typing, it does nothing to address the excessive redundancy of data inherent in such structures, which, when such trees are traversed at session initiation time, burdens consumer and providers alike.

The purpose of templates described in this extension is to provide a way for providers to largely reduce these redundancies by means of defining templates holding a common set of parameters, attributes and even sub-trees as parts of the Ember+ tree, which then can be referred to by other elements, allowing the provider to only communicate differences to the referred template element in descriptive messages or value changed notifications.

Specification within an Ember+ tree

Template definitions

Template definitions may occur anywhere within an Ember+ tree, provided that the context, in which a template definition appears is also a valid context for the type of element it describes. Being an element in the Ember+ nomenclature, a template may describe the structure of a parameter, a node, a matrix, or a function. It may not describe a command, since commands are a fixed set of operations issued by a consumer and do not serve descriptive purposes within a tree offered by the provider.

Scope

The scope of a template definition spans the entire sub-tree spanned by the immediate root node, below which the template definition is encountered and does not extend to other root nodes nested within this sub-tree.

Template references

An element may refer to a template definition by means of the newly introduced `templateReference` attribute that is part of the element's `Contents`. The value of this attribute must be a relative OID relative to the immediate root node common to the template definition and the element referring to it.

Restrictions, semantics and precedence

Template identity

The number of a template is the number of its contained element definition. While it might be convenient to consider a reference to a template as a “copying the template definition into the referring node”, this interpretation is incorrect, which results in the constraint that a template definition nested within another template definition exists only once in a tree and may only be referred to by the path leading to its original, explicit definition.

Acyclicity

The contents of a template definition can be any definition of the abovementioned element types, which may also include nested template definitions as well as references to other templates, with the only restriction being, that the directed template dependency graph of a template definition must be acyclic. For the sake of this cyclicity check a template is considered defined, once its Template element has been encountered, even though its Number, which gives the template a referable label, is actually nested within.

Type equality

The types of the referred template and the referring element must be the same.

Constness

A template definition has to remain constant over the course of a session. That means, that:

1. Subsequent traversals of the template definition must not report differing contents.
2. No value change notifications must be emitted for template elements.

On the other hand a template reference reported by a node must not change over the course of a session.

Semantics

When an element refers to a template definition it subsumes the contents of the referred template element with the exception of nested template definitions, as mentioned above. For an element governed by a template reference¹ all attributes are optional as long as all attributes that are declared as required by the DTD are contained in the union of the provided as part of the template definition and the attributes explicitly provided by the element.

If the provider explicitly enumerates an element, governed by a template definition, the attribute Number is required.

¹ An element is said to be “governed by a template reference” if it either contains a templateReference attribute, or one of its parent nodes refers to a template that defines a child element at the same location.

Precedence

When an element enumerates an attribute or child element during traversal or sends a value changed notification over the course of a session, the reported values take precedence over the corresponding values contained in the template definition.

The attributes or child elements enumerated or changed by an element need not be a subset of the contents of the template.

GetDirectory requests

It is at the direction of the consumer whether or not it wants to utilize the bandwidth-saving capabilities of template definitions and references. Therefore the semantics of the already defined explicit `dirFieldMask` values remain unchanged. Most importantly a provider must enumerate all contents of an element, regardless of their redundancy with regard to a referred template, when the consumer requests All (-1) element contents. This extension introduces a new `dirFieldMask` value that explicitly requests non-redundant content enumeration. It is named Sparse and is signified by the value -2.

Compatibility and migration

This extension will introduce a new version of the Glow DTD, which is why compatibility for this extension must be indicated by reporting the corresponding version in the `AppBytes` field of the initial S101 frame. A provider may assume this version to be constant and disregard version changes in later frames.

Provider migration

Migration of a provider is simple, as the semantics of tree elements, as well as the `GetDirectory` command remain unchanged if a tree does not contain template definitions or template references. In this case the new `dirFieldMask` value Sparse is equivalent to All, which also implies unchanged semantics of the value Default.

Consumer migration

A consumer that is not yet fully compatible with templates can implement a minimal support strategy that makes it compatible without any major changes to its internal model. Since the `dirFieldMask` value All retains its old semantics and is equivalent to the old Default value. So all a consumer has to do is to:

1. Always use All as the `dirFieldMask` in `GetDirectory` requests.
2. Discard all Template definitions or references encountered as part of a response to a `GetDirectory` request.

Changes to the Ember+ DTD

The following two types will be added to the DTD:

```

Template ::=
    [APPLICATION 24] IMPLICIT SET {
        Number          [0] Integer32
        Element          [1] TemplateElement
        Description      [2] EmberString
    }

QualifiedTemplate ::=
    [APPLICATION 25] IMPLICIT SET {
        Path             [0] RELATIVE-OID
        Element          [1] TemplateElement
        Description      [2] EmberString
    }

TemplateElement ::=
    CHOICE {
        parameter    Parameter,
        node         Node,
        matrix       Matrix,
        function     Function
    }

```

The types listed below contain changes:

```

Element ::=
    CHOICE {
        parameter    Parameter,
        node         Node,
        command      Command,
        matrix       Matrix,
        function     Function,
        template     Template
    }

RootElement ::=
    CHOICE {
        element             Element,
        qualifiedParameter  QualifiedParameter,
        qualifiedNode       QualifiedNode,
        qualifiedMatrix     QualifiedMatrix,
        qualifiedFunction    QualifiedFunction,
        qualifiedTemplate    QualifiedTemplate
    }

```

```

ParameterContents ::=
  SET {
    -- Attributes ...
    schemaIdentifiers    [17] EmberString    OPTIONAL,
    templateReference    [18] RELATIVE-OID    OPTIONAL
  }

```

```

NodeContents ::=
  SET {
    -- Attributes ...
    schemaIdentifiers    [4]  EmberString    OPTIONAL,
    templateReference    [5]  RELATIVE-OID    OPTIONAL
  }

```

```

MatrixContents ::=
  SET {
    -- Attributes ...
    schemaIdentifiers    [11] EmberString    OPTIONAL,
    templateReference    [12] RELATIVE-OID    OPTIONAL
  }

```

```

FunctionContents ::=
  SET {
    -- Attributes ...
    result                [3]  TupleDescription OPTIONAL,
    templateReference    [4]  RELATIVE-OID    OPTIONAL
  }

```

```

FieldFlags ::=
  INTEGER {
    sparse      (-2),
    all         (-1),
    default     ( 0),
    identifier   ( 1),
    description  ( 2),
    tree        ( 3),
    value       ( 4),
    connections ( 5),
  }

```

Ember+ usage

With the object types defined in the previous chapter it is now possible to represent a piece of hardware, for example a frame controller which has one or two power supplies and ten slots for different kinds of cards. All types mentioned before are available in the Ember library, what makes it intuitive and easy to use. This chapter describes how devices are being represented in Glow.

Representing a device

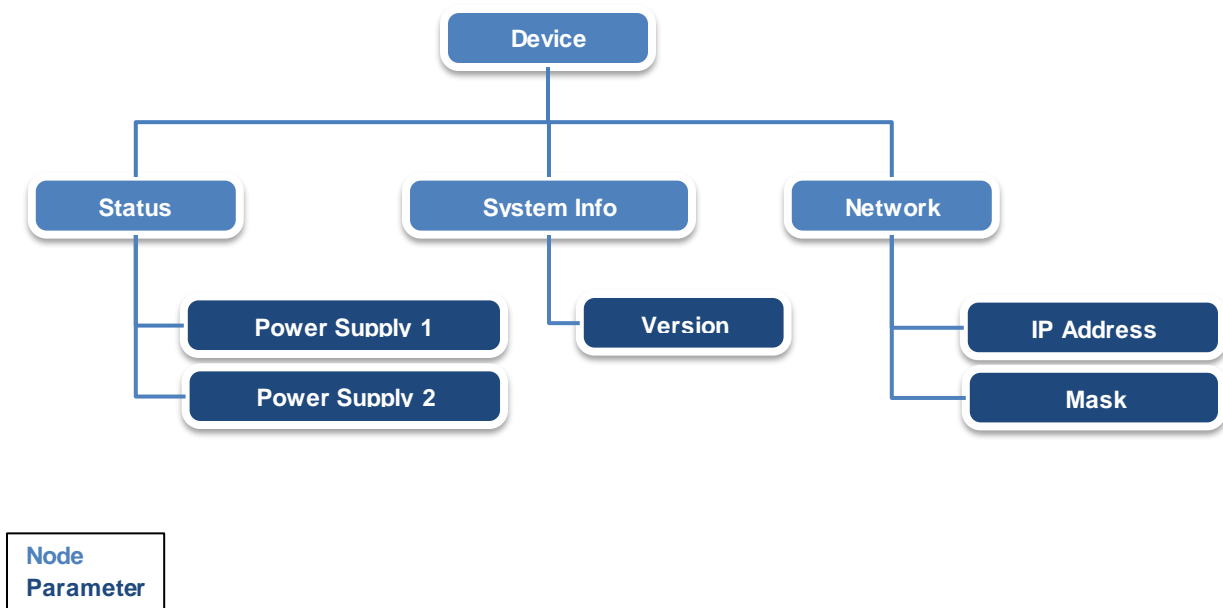
A common device is represented by using the Node and the Parameter type. Each node may have any number of child nodes and parameters, while a parameter may not have child nodes or parameters – although the specification technically allows it since the Children property is of type ElementCollection. But this property should only be used by a consumer to attach a command to the parameter, like [Subscribe](#) or [Unsubscribe](#).

The sample frame provides several parameters, which are listed in the table below.

Identifier	Type	Properties
Power Supply 1	Integer (Enumeration)	<ul style="list-style-type: none">• Value• Enumeration• Description• Identifier
Power Supply 2	Integer (Enumeration)	<ul style="list-style-type: none">• Value• Enumeration• Description• Identifier
Software Version	String	<ul style="list-style-type: none">• Value• Description• Identifier
IP Address	String	<ul style="list-style-type: none">• Read/Write• Value• Description• Identifier
Network Mask	String	<ul style="list-style-type: none">• Read/Write• Value• Description• Identifier

This is a simplification and only a subset of the parameters a device usually offers, but it is enough to show how this structure would look like when using Glow. The easiest approach would be to create a Node representing the device and then to append all parameters to it. This

would work but the Node type is not restricted to representing physical entities, it may also be used for logical entities, like a category.



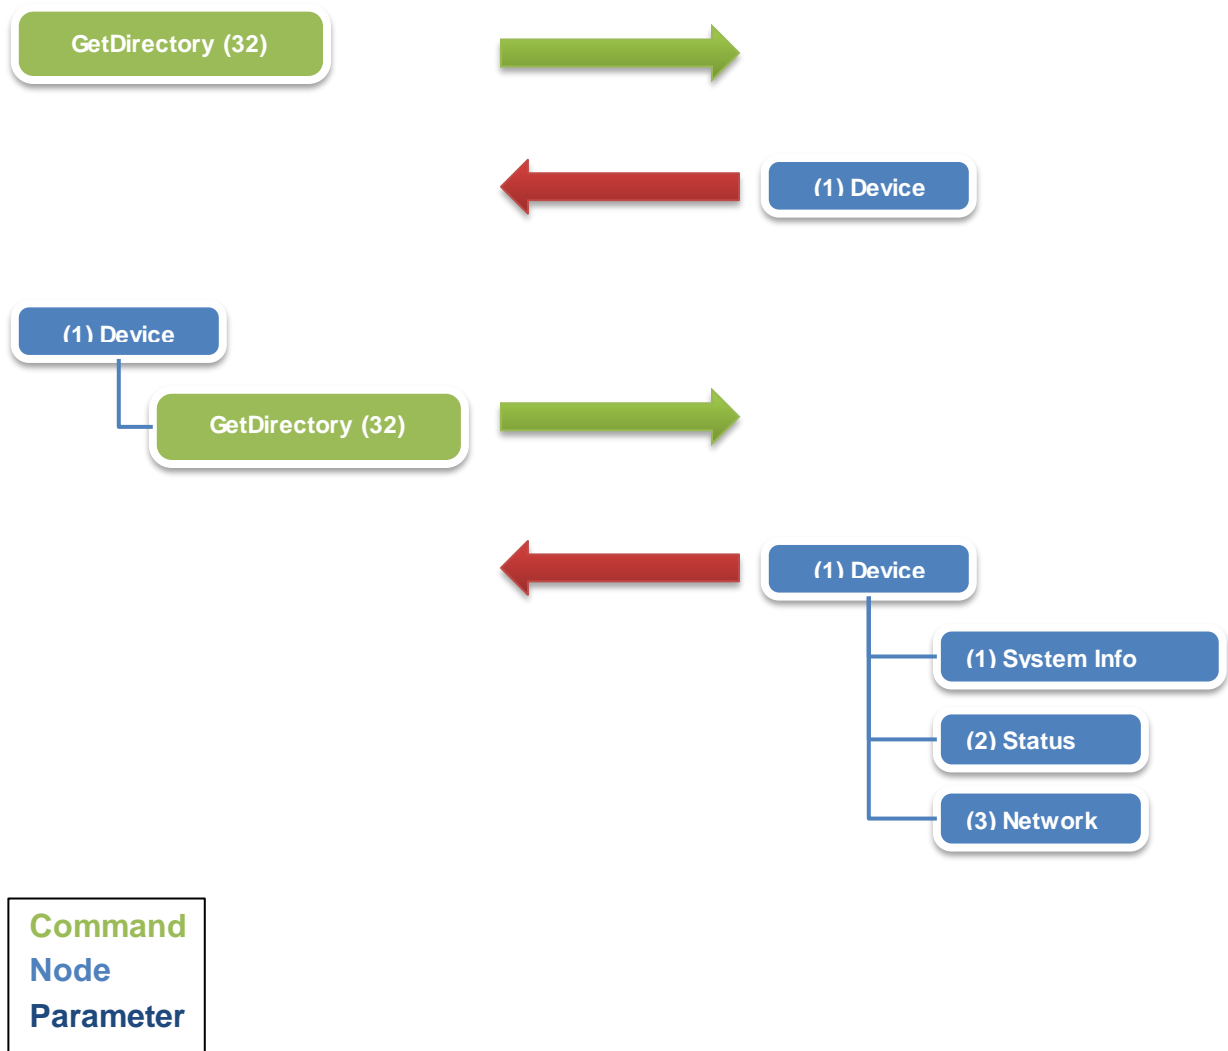
This design separates the parameters into three categories: *Status*, *System Info* and *Network*. These names are also the node identifiers. In most cases, the benefit is that the user will find the parameter it is looking for faster. A consumer may also use this information for the layout of its user interface. In bigger systems with hundreds or even thousands of parameters the use of nodes in order to categorize parameters also reduces the message sizes and reduces the CPU load of the data provider since it doesn't have to encode a large parameter set at once. Instead, it only has to transmit the parameters of the node the user or a control system is currently interested in.

Querying a data provider

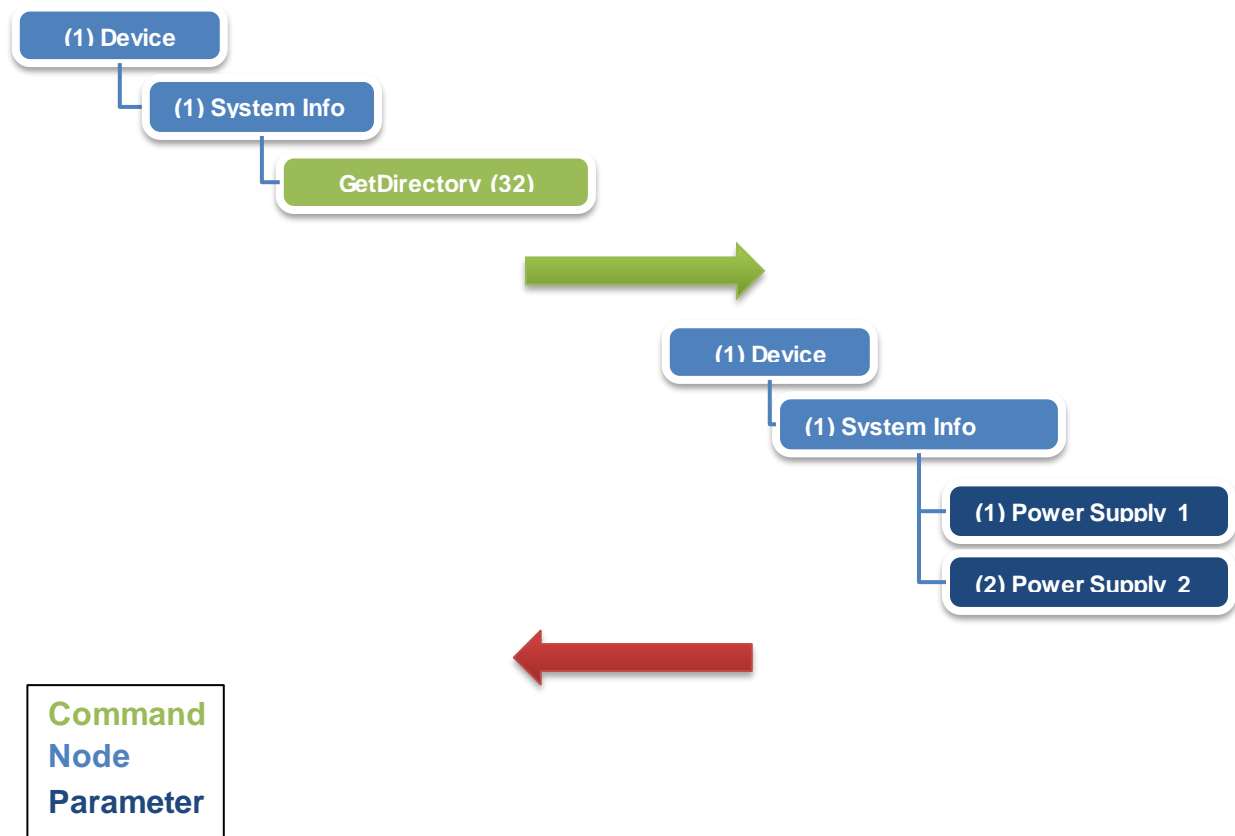
When a consumer connects to a provider it usually requests the structure of the device. This can be done by using the [GetDirectory](#) command, which has to be sent for every node the consumer is interested in.

Basic data exchange

The initial data exchange would look like this:



The consumer sends a `GetDirectory` request (which must be the child element of the root element collection) to the provider, which responds with another `ElementCollection` containing the root node “Device” with Number 1 (written within the parentheses) and its description. To query the child elements of the device node, the consumer must create a message that contains the device node and append a `GetDirectory` command to the `Children` property of that node. The provider then responds with the device node and all of its children. When the consumer is interested in the child elements of the “System Info”, it would send the following request:



The consumer appends the “GetDirectory” command to the node it wants to get the children from, in this case “System Info”. The data provider then responds with the same structure, but appends the child nodes and parameters, if there are any. When a provider reports parameters due to a “GetDirectory” request, it must include all relevant properties. That way a consumer can query the complete structure of any data provider.

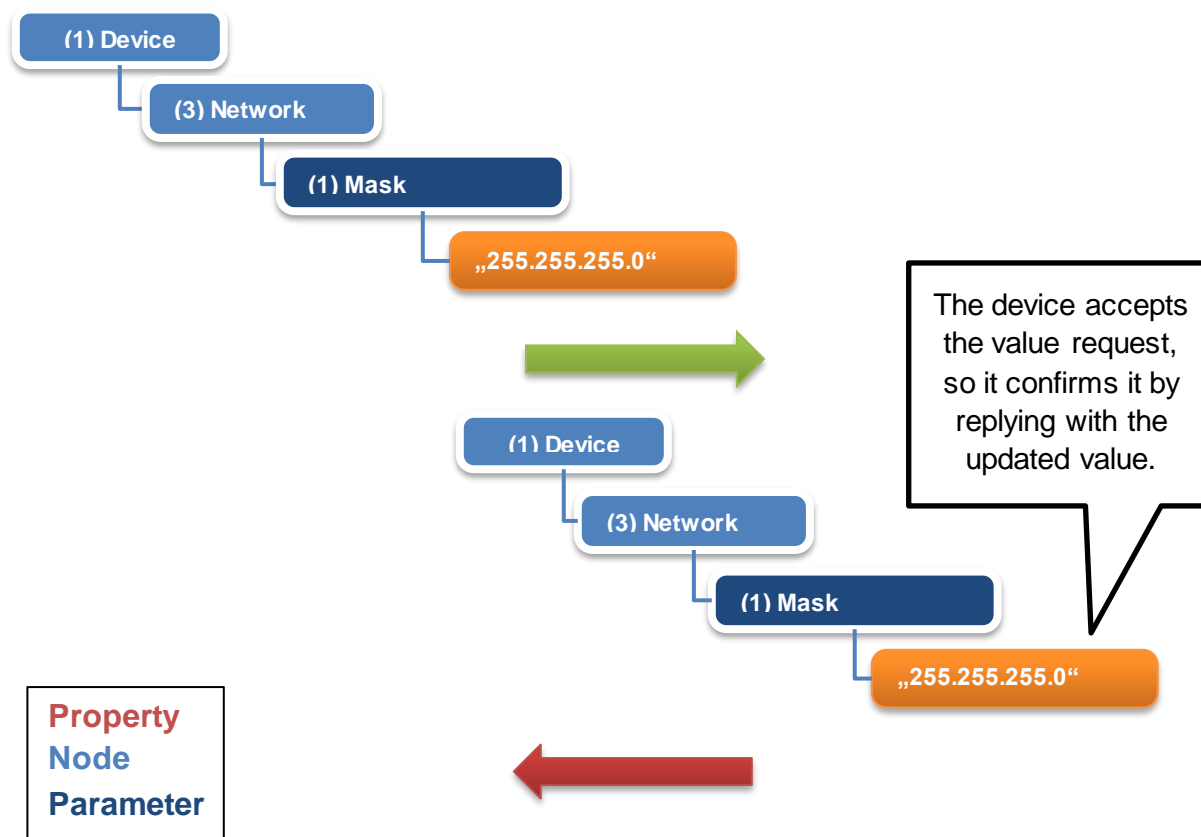
Changing a parameter value

When a consumer has received the structure of a provider, a user or a control-system now has the possibility to modify the writeable values of the parameters available. There is no additional command required in order to change a parameter. Basically, the consumer only has to send the value that should be changed to the provider. The small sample structure used before has two writeable parameters: IP Address and Network Mask.

Please note that the consumer could lose its connection to the provider when it changes the IP Address and uses TCP/IP or UDP as transport layer. This example is for demonstration purposes.

The consumer must transmit the structure containing the parameter it wants to change, and finally the value to set the parameter to. When a provider receives a value change request, it

must evaluate the new value. If it is valid the provider must apply the requested change and respond with the new value. Otherwise, if the value received is invalid for some reason, the provider should discard the request and respond with the current value, even if it doesn't change. The following sample shows how a value change request looks like.



Note: In the examples above, the messages from the provider contain only normal elements (e.g. nodes and parameters). A provider is however free to instead send messages with the qualified variants of these elements or even a combination of normal and qualified elements. Vice versa, requests issued by the consumer can also contain any combination of normal and qualified elements.

The EmBER library

The implementation of the EmBER library is written in standard C++ 03 and has been tested on several different compilers (MSVC++, g++) and platforms (Windows, Mac OS X). This section describes the different layers of the library and how to use it.

The library has been designed to work on most embedded systems as well.

Compiling the library

The library can be built with any C++ 03 standard compliant compiler. There are two variants, it may either be compiled or used as library, or it can be used header-only, which means that an application only includes the necessary headers.

To use the library as header-only, define the “LIBEMBER_HEADER_ONLY” macro before including the header file(s):

```
#define LIBEMBER_HEADER_ONLY
#include "ember/Ember.hpp"

// ...
```

A premake file is provided with the library which can be used to compile the source code when not using the header-only variant. For more information, please visit:

<http://industriousone.com/premake>

Including the library

Each layer/namespace provides a convenience header, which should be used instead of including the files separately. If all layers are being used, the file “Ember.hpp” can be included. Otherwise, the files listed in the table below can be used:

Namespace	File
ber	“ber/Ber.hpp”
dom	“dom/Dom.hpp”
glow	“glow/Glow.hpp”

Library structure

The library has been divided into four different layers, as listed in the table below.

Util

The util namespace contains several helper classes and macros, especially the TypeErasedIterator and the OctetStream.

BER

The BER layer is the low level layer which provides the basic encoding and decoding routines for all data types supported, including the length type and the tag. The implementation uses traits classes in order to provide the encoding and decoding algorithms. That way it is easy to extend the library by additional data types.

Dom

The Dom layer provides the classes that are used to build an encodable data structure.

Glow

The Glow layer is based on the Dom layer and provides classes for each object type defined by the Glow specification, like `GlowNode`, `GlowParameter`, `GlowCommand` and `GlowElementCollection`.

The util namespace

The util namespace contains some macros and helper classes which are used all over the library, like the `TypeErasedIterator` and the `StreamBuffer`. The iterator class is used by the container classes implemented in the dom and glow namespace and allows to traverse the children of a container. The `StreamBuffer` is used by the encoding and decoding functions. The `StreamBuffer` uses a singly linked list where each node allocates a fixed size chunk.

The BER namespace

As mentioned in the introduction, this namespace contains the encoding and decoding functions. It also provides classes for the Tag, the Length and a generic Value.

To encode a value, for example an integer, its universal tag and the encoded length must be known. The universal tag contains the type information, for an integer this would be “Universal - 2”. All universal types are enumerated in the `Type` class (see *ember/ber/Type.h*).

The following example demonstrates how a value is being encoded.

```
util::OctetStream stream;
ber::Tag const appTag = ber::make_tag(ber::Class::Application, 1);

int const intValue = 1333;
std::size_t const frameLength = ber::encodedFrameLength(intValue);

ber::encode(stream, appTag);
ber::encode(stream, ber::make_length(frameLength));
ber::encodeFrame(stream, intValue);
```

The output buffer contains the following encoded bytes:

0x41 Application Tag “Application-1”**0x04 Inner Frame Length, 4 Bytes****0x02 Type Tag “Universal-2” (integer)****0x02 Data Length, 2 Bytes****0x05 0x35 Encoded value, 1333**

To decode this data, the following sample can be used:

```
appTag = ber::decode<ber::Tag>(stream);  
frameLength = ber::decode<ber::Length<int>>(stream);  
universalTag = ber::decode<ber::Tag>(stream);  
length = ber::decode<ber::Length<int>>(stream).value;  
intValue = ber::decode<int>(stream, length);
```

This example is very basic and only demonstrates the basic use of the library and how the encoded result looks like.

In a real world example, it would be better to use the `DecoderFactory` class instead. It decodes a set of type, length and value and returns an instance of the type-erased `Value` class.

The Dom namespace

The Dom namespace offers a set of classes that can be used to create an encodable structure. A programmer using these classes does not have to have detailed knowledge about the encoding and decoding rules.

The namespace also contains two classes called `DomReader` and `NodeFactory`. The `DomReader` is used to decode a buffer which contains a complete tree. Additionally, it may be used to decode data which uses application defined types, like the ones used in the Glow specification. This is why the reader requires an instance of a `NodeFactory`, which is an abstract class responsible for creating concrete objects. It is always being called by the reader when it detects an application defined object.

The Glow namespace

The classes inside this namespace are built upon the dom classes. They represent the object types used by the Glow specification. The classes within this namespace allow it to easily construct a valid Glow message and they can also be used to traverse a decoded ember message.

Using the glow classes

This section briefly demonstrates the use of the Glow classes by providing some simple examples.

Creating a GetDirectory request

The first example shows how to generate a GetDirectory request at root level.

```
root = GlowRootElementCollection::create();

// The create method returns a collection with the application tag
// set to GlowTags::Root.

new GlowCommand(root, CommandType::GetDirectory);

// Encode and transmit
// ...
```

The next example shows how to query the children of a specific node. It uses the small example from the previous chapters. This snippet queries the children of the Network node. The complete path up to the node that is being queried must be provided, but it is sufficient to add the node numbers. The last node then contains the GetDirectory command as child element.

```
root = new GlowRootElementCollection::create();
device = new GlowNode(1);
network = new GlowNode(3);

command = new GlowCommand(CommandType::GetDirectory);

root->insert(root->end(), device);
device->children()->insert(device->children()->end(), network);
network->children()->insert(network->children()->end(), command);
```

It is also possible to query several nodes within one message. A consumer could construct a single tree with all known nodes and append a GetDirectory command to each of them.

Replying to a GetDirectory request

When a provider receives a valid GetDirectory request it must respond with all nodes that belong to the one being queried. In our example, the network node contains two parameters, IP Address and Mask. So the response to the previous query would look like this:

```
root = GlowRootElementCollection::create();
device = new GlowNode(1);
network = new GlowNode(3);

ipAddress = new GlowParameter(1);
ipAddress->setValue(currentIpAddress);
ipAddress->setAccess(Access::ReadWrite);
ipAddress->setIdentifier("ipaddr");
ipAddress->setDescription("IP Address");
```



```

mask = new GlowParameter(2);
mask->setValue(currentMask);
mask->setAccess(Access::ReadWrite);
mask->setIdentifier("netmask");
mask->setDescription("Network Mask");

root->insert(root->end(), device);
device->children()->insert(device->children()->end(), network);
network->children()->insert(network->children()->end(), ipAddress);
network->children()->insert(network->children()->end(), mask);

```

This way a message is being created which contains the current state of the parameters and their description.

Changing a parameter value

When a consumer wants to change a parameter value, it has to provide the partial tree and the new value. The following sample shows how a change network mask request looks like.

```

root = GlowRootElementCollection::create();
device = new GlowNode(1);
network = new GlowNode(3);

mask = new GlowParameter(2);
mask->setValue("255.255.252.0");

root->insert(root->end(), device);
device->children()->insert(device->children()->end(), network);
network->children()->insert(network->children()->end(), mask);

```

This is it. When a provider detects a parameter containing a value, it must treat this value as value change request.

Reporting a parameter value change

When a provider receives a value change request or changes a value for some other reason, it must report the change to all connected clients. The response to the previous example looks like this:

```

root = GlowRootElementCollection::create();
device = new GlowNode(1);
network = new GlowNode(3);

mask = new GlowParameter(2);
mask->setValue(currentMask);

root->insert(root->end(), device);
device->children()->insert(device->children()->end(), network);
network->children()->insert(network->children()->end(), mask);

```

When the consumer transmits a valid value, the response should always look like the request. If the value provided by the consumer is invalid, the provider answers with the current value.

Traversing a tree

Both, the consumer and the provider must traverse a tree when they receive it. In both cases, the procedure is very similar. This section demonstrates how a provider may traverse a tree. The sample covers Nodes, Commands and Parameters.

```
void traverseChildren(GlowElementCollection& collection)
{
    GlowElementCollection::iterator it = collection.begin();
    GlowElementCollection::iterator last = collection.end();

    for(; it != last; ++it)
    {
        dom::Node& node = *it;

        // Converts the node's type tag into a ber type. That makes it easier
        // to determine if this is an application defined type.
        ber::Type type = ber::Type::fromTag(node.typeTag());

        if(type.isApplicationDefined())
        {
            switch(type.value())
            {
                case GlowType::Node:
                    handleNode(dynamic_cast<GlowNode&>(node));
                    break;
                case GlowType::Parameter:
                    handleParameter(dynamic_cast<GlowParameter&>(node));
                    break;
                case GlowType::Command:
                    handleCommand(dynamic_cast<GlowCommand&>(node));
                    break;
            }
        }
    }
}

void handleNode(GlowNode& node)
{
    // An implementation should validate if the current node
    // really exists by checking its number.

    // Then, traverse the children - if there are any.
    GlowElementCollection* children = node.children();

    if(children != 0)
        traverse(*children);
}

void handleParameter(GlowParameter& parameter)
```

```
{
    // If this is a provider side implementation, the
    // provider usually only has to check if a value is set.
    // If so, this is a value request that must be handled:

    if(parameter.contains(GlowProperty::Value))
    {
        // Handle Set-Value request:
        handleParameterValue(parameter, parameter.value());
    }

    // On the other side, if this is a consumer implementation, it should check
    // the availability of each property and update its internal state.

    if(parameter.contains(GlowProperty::Access))
    {
        // modifyAccessparameter, parameter.access()
    }

    if(parameter.contains(GlowProperty::Minimum))
    {
        // modifyMinimum parameter, parameter.minimum()
    }

    // and so on . . .
}
```

Behaviour rules and other guidelines

This chapter lists the general rules a device must follow when it uses the EmBER protocol. Additionally, this chapter provides some samples that demonstrate how a tree should look like – or how it shouldn't look like.

General behaviour rules

Message length

To avoid packets that may be too large for a provider or consumer with lower memory capabilities, the payload size of an Ember+ packet is limited to 1024 bytes. Due to framing overhead, the actual maximum length of an Ember+ packet is 1290 bytes. The Ember+ libraries provide methods to transmit partially encoded data. The framing protocol, which is used to transport the encoded EmBER data, provides a set of flags which are being used to indicate whether the encoded data is only a part of a message or a complete packet. More details about these flags and S101 in general can be found under [S101 Messages](#).

The *StreamBuffer* class from the C++ library for example provides a virtual flush method, which is invoked whenever the buffer reaches its maximum size, which can be determined via a template parameter. When overriding this method, it is possible to encode transmit the current data before the buffer content is being reset.

That way, the memory usage is kept low and the data is being written to the clients while the EmBER encoding is in process.

On the side, when receiving a partial message that needs to be decoded, the library provides the *AsyncDomReader* class, which invokes the virtual *rootReady* method as soon as the decoding is complete. It is possible to derive from that class and override this method in order to process the decoded tree.

The identifier property

The identifier of a node must never change. Otherwise a consumer cannot refer to a specific node or parameter any more. If, for example, a device has several slots (1-10) which may contain different kinds of modules, there must be one unique identifier for each slot – module combination. So a node called “Slot_01” would not satisfy this requirement. The identifier should additionally contain some module information, like a type identifier or a version. The identifier could then look like “Slot_01_1000_2.0”.

The identifier string must not contain the character “/” (ASCII/UTF-8: 47) and must begin with a letter or the lower line character (“a”-“z”, “A”-“Z”, “_”, ASCII/UTF-8: 65-90, 97-122 ,95).

The description property

If a node or parameter provides a description property, it shall be displayed instead of the identifier. The identifier is not required to be human readable in a means that it has to make sense for a user.

The number property

The number is used to quickly identify a node or parameter. Once a consumer knows the structure of a data provider, it should always only provide the node or parameter numbers instead of the string identifier.

When to use the element number or identifier

Both the provider and the consumer must always include the element (node or parameter) number when transmitting a tree.

When a consumer queries the nodes of a provider by using the [GetDirectory](#) command, it must provide all attributes of the child elements, including the string identifier.

The consumer itself is never required to transmit the identifier. But if it receives an element which contains both, a number and an identifier, the identifier shall be used to map the node to an existing internal structure. This must usually only be done once when the consumer queries the provider for the first time after it established a connection.

Keep-Alive mechanism

A basic Keep-Alive mechanism can be implemented by using the S101 commands called “Keep-Alive Request” and “Keep-Alive Response”. Both, the provider and the consumer may use the “Keep-Alive Request” command in order to test whether the remote host is still responsive.

Though issuing the “Keep-Alive Request” message is optional, any Ember+ host **MUST** respond to a “Keep-Alive Request” message with a “Keep-Alive Response”.

The host sending the “Request-Keep Alive” message determines how long it waits for a response. A recommended timeout is about 5 seconds. The host that issued the keep-alive request may close the connection to the remote machine when the timeout interval expires. The keep-alive request should only be sent when there hasn't been any kind of communication for a while, like for about 4-5 seconds. For more information, please refer to the chapter [S101 Messages](#), which describes the S101 commands.

Number of consumers per provider

A provider should not limit the number of active client connections to one or two. Most modern control systems run in redundant mode and therefore it may be required that there are several connections active at once. This of course depends on the hardware capabilities of the provider, but it should be considered when designing the remote interface.

Notifications

Often a device allows several ways to be controlled, not only the remote interface. A second way would be a front panel which allows modifying some parameters. To avoid polling from the

consumer side(s), a provider must always automatically notify its clients about the parameter change. It is only required to inform the clients that have already sent a GetDirectory request to the parent node of the affected parameter. This includes the online/offline state of a node and all properties of a node or parameter. These notifications must not be sent when the parameter has a stream identifier, which indicates that it must be explicitly subscribed.

When a parameter value changes, the provider must at least transmit the new value. All other parameter attributes must not necessarily be transmitted. However, they may be transmitted. The change notifications must be transmitted to all clients that are currently connected.

Parameter value range changes

Although it is possible to change the value range of a parameter it is not recommended. This usually irritates users working with a user interface, especially when the range changes every time when a second parameter has been modified.

Value change requests

A provider must always respond to a value change request. If the value is invalid, the provider shall respond with the current value.

Remarks on the Behaviour of Embedded Devices

If an embedded device with very low hardware profile (well below ARM processor) wants to act as an Ember+ provider, it is allowed to transmit its entire static Glow tree once as soon as a client connects – without processing the GetDirectory command at all. So every time the provider accepts a connection from a new consumer, it sends its entire tree to the consumer and from then on only processes parameter updates sent from the consumer.

Thus, the binary representation of the static tree may be compiled into the embedded software. This is only allowed for providers which publish a maximum tree depth of 2 and a maximum of 32 parameters. These static trees must not contain empty nodes.

Removing Sub-Trees at Runtime

When a provider wants to notify consumers that a certain sub-tree no longer exists (e.g. if a module was removed at runtime), it should send a notification to all consumers containing the parent node of the removed sub-tree with the property “IsOnline” set to false.

From then on the provider no longer includes the removed sub-tree in responses to the GetDirectory command.

General recommendations

Separate parameters into logical categories

To avoid very large parameter lists it is recommended to introduce logical categories, like “Audio”, “Video”, “Settings”.

Avoid indexed parameters

A router for example often provides a gain parameter for each source or target it has. It is NOT recommended to visualize this by providing a single node which contains all gain parameters with their number in their name. Instead there should be one node for each target or source which has the name of the target ("Target 1") and then contains a parameter called "Gain". The table below demonstrates both variants:

Don't do this	Recommended variant
<ul style="list-style-type: none">▪ Target Gains<ul style="list-style-type: none">▪ Gain [0]▪ Gain [1]▪ Gain [2]▪ ...▪ Gain [N]	<ul style="list-style-type: none">▪ Targets<ul style="list-style-type: none">▪ Target 1<ul style="list-style-type: none">▪ Gain▪ Target 2<ul style="list-style-type: none">▪ Gain▪ ...

Interval for stream transmission

The recommended interval to transmit streams a consumer has subscribed to is between 50 and 80 milliseconds.

Offline trees

In several scenarios it is a requirement to store the tree structure of a provider on the consumer side. This might be the case, when a device is not accessible at any time, but its tree is required for some configuration tasks by a consumer (like a control system).

For this situation, it is possible to store the device's tree by querying all data that is required once and then storing it by using an Ember+ encoder or any user defined data format. The tree can then be accessed by the consumer as if the device was connected. It would even be possible to load the data into a kind of Ember+ emulator software. This would allow simulation of value control.

But it is important to note that the node numbers and parameters may have changed when reconnecting to the device at a later point of time. So it is important for a consumer to reference nodes by using their string identifiers instead of their number when storing a tree structure for offline use.

Frequently asked questions

This section contains a set of questions that came up during the development of the library and the integration into other projects.

Why are the node and parameter numbers session based?

The majority of the providers use a static numbering scheme for its nodes and parameters. This is the case when the hardware is known and doesn't change. But devices with modular equipment need to dynamically generate node numbers at runtime, for example when a card is being replaced by another card of a different type.

Message Framing

In order to transport the EmBER encoded data, an additional framing is required which allows the packets to be transmitted via any transport layer. This task is done by the S101 protocol, which uses a start and end byte to indicate transmission begin and end and a 16 Bit CRC which can be used to validate the packet.

The S101 protocol

To assure that the start or end byte of the framing doesn't appear within a message, all bytes with a value above 0xF8 are being escaped with a special character. First of all, the following table shows all special characters used in S101.

Name	Value	Description
BOF	0xFE	Begin of Frame
EOF	0xFF	End of Frame
CE	0xFD	Character escape
XOR	0x20	XOR value for Character escape
Invalid	0xF8	All bytes above or equal to this value must be escaped

Encoding a message

When encoding data, each data byte must be compared against the S101 Invalid value. If the value is below 0xF8, the byte can be appended to the encoding buffer. Otherwise, it must be escaped. This is done by adding the CE (0xFD) character into the stream followed by the original byte XOR'ed with XOR (0x20). Additionally each message must start with the BOF character and end with the 16 Bit CRC and the EOF character.

The data 0xFF, 0x00, 0xF9, 0x01 would encode to:

```
0xFE, 0xFD, 0xDF, 0x00, 0xFD, 0xD9, 0x01, 0x95, 0x83, 0xFF
```

Because 0xFF is above the Invalid character, it will be escaped which results in 0xFD, 0xDF. The same mechanism is applied to 0xF9. The algorithm to compute the CRC is described at the end of this chapter.

Decoding a message

When a decoder reads a BOF byte, it always indicates the start of a new frame. So the current data can be discarded if there is any. When the decoder receives a CE character, the current byte can be discarded and the next byte must be XOR'ed with 0x20 and then added to the receive buffer. When the decoder reads the EOF, the message is complete and the checksum can be evaluated.

CRC Computation

All data bytes are used for the CRC computation. The result is inverted (by applying the binary NOT operator) and the stored after the data bytes and before the EOF. The CRC bytes must also be escaped when they are equal to or above the S101 Invalid byte.

The initial value of the CRC is always 0xFFFF. The result of a decoded CRC must always be 0xF0B8. The CRC table can be found in the appendix. The following function must be applied to each data byte that is being encoded:

```
static const WORD table[256] = { . . . };

WORD ComputeCRC(WORD crc, BYTE byte)
{
    return (WORD)((crc >> 8) ^ table[(BYTE)(crc ^ byte)]);
}
```

S101 Messages

The content of all S101 messages starts with the slot identifier, which in this case is usually set to 0x00. The second byte contains the message type, which is set 0x0E for Ember. The third byte determines the command. The meaning of all other appended bytes depends on the command.

Message types

0x0E	EmBER
<i>The range 0x00 to 0x10 is reserved for internal use.</i>	

Commands

0x00	EmBER Packet
0x01	Keep Alive Request
0x02	Keep Alive Response

Messages

This section describes the format of the available messages.

EmBER Packet

The EmBER packet command is used whenever a device transmits data encoded with BER. The packet has the following format:

Slot	Message	Command	Version	Flags	DTD	App Bytes	<Payload>
------	---------	---------	---------	-------	-----	-----------	-----------

Slot

The slot is used to address a kind of sub-device within a device. Right now, this byte should be set to 0x00.

Message

The message type; set to 0x0E.

Command

The command determines the content of the following message bytes. Available commands are EmBER (0x00), Keep-Alive Request (0x01), Keep-Alive Response (0x02).

When the command is EmBER (0x00), the payload contains an encoded Ember packet. When the command is Keep-Alive Request (0x01), the client (which may be provider or consumer) must response with a Keep-Alive Response (0x02).

Version

Set to 1.

The following fields are only available when the command type is EmBER (0x00). For Keep-Alive requests and responses, only Slot, Message type, Command type and Version must be transmitted.

Flags

The upper three bits of the flags are used to indicate partitioned ember packets. The following table lists the flags and their meaning:

0xC0	Single-packet message
0x80	First multi-packet message
0x40	Last multi-packet message
0x20	Empty packet
0x00	A packet within a multi-packet message

DTD

The DTD byte defines the “Design Type Document” that is being used. This value must be set to one, which means that the Glow specification is used.

App Bytes

This field contains the number of bytes that follow before the payload begins.

The current version uses two application bytes, which contain the version number of the Glow DTD.

Payload

The Payload contains the BER-encoded data. This data may be decoded via a BER decoder or the DomReader.

Usage

The basic header of a non-partitioned single packet message would look like this:

```

0x00 Device Slot
0x0E Message type: EmBER
0x00 Command type: EmBER Packet
0x01 Version
0xC0 Flags: Single Packet message (First packet | Last packet)
0x01 DTD Type: Glow DTD
0x02 Application Bytes
0x05 Minor Glow DTD Version
0x02 Major Glow DTD Version
[EmBER Data]
```

Glow DTD ASN.1 Notation

This section contains the ASN.1 notation of the Glow specification.

```
--
-- GlowDtd.asn1
-- Lawo GmbH
--
-- This file defines the Glow DTD used with the Ember+
protocol.
--
-- Change Log:
--
-- 2.40:
--   - NOTE: This version describes the data schema of Ember+
1.4.
--   - Added Template Extensions
-- 2.30:
--   - NOTE: This version describes the data schema of Ember+
1.3.
--   - Added Schema Extensions
-- 2.20:
--   - NOTE: This version describes the data schema of Ember+
1.2.
--   - Added Function Extensions (see type Function)
-- 2.10:
--   - NOTE: This version describes the data schema of Ember+
1.1.
--   - Added Matrix Extensions (see type Matrix)
--   - Added "isOnline" field to NodeContents
-- 2.5:
--   - NOTE: This version describes the data schema of Ember+
1.0.
--   - NOTE: This version introduces breaking changes!
--   - Changed Parameter.isCommand (BOOLEAN) to an
enumeration named "type".
--   To determine the effective type of a parameter, follow
this rule:
--   - If the parameter has the "type" field that equals
"trigger", its
--     type is "trigger".
--   - If the parameter has either the "enumeration" or the
"enumMap" field,
```

```
--      its type is "enum".
--      - If the parameter has the "value" field, its type
corresponds to the
--      BER type of the value.
--      - If the parameter has the "type" field, its type is
the value of this
--      field.
--      This is useful for parameters that do not specify a
current value -
--      e.g. "trigger" parameters or parameters that have
write-only access.
--      - Changed Parameter.isWriteable (BOOLEAN) to an
enumeration named
--      "access".
--      - More options for Value - now also supports OCTET
STRING and BOOLEAN
--      - Introduces QualifiedParameter and QualifiedNode types
--      - Introduces RootElement and RootElementCollection
types:
--      At the root level, a different set of supported types
is available.
--      - StreamCollection can also be used as root container.
--      - Introduces the StreamDescription type and the field
"streamDescriptor"
--      in type ParameterContents.
-- 2.4:
--      - NOTE: This version introduces breaking changes!
--      - moved "children" in Parameter and Node out of
--      "contents" SET.
-- 2.3:
--      - Added size constraints for INTEGER values.
--      - Renamed EnumEntry to StringIntegerPair
--      - Renamed EnumCollection to StringIntegerCollection
-- 2.2:
--      - Added new field "enumMap" to Parameter and types to
describe
--      enum entries: EnumEntry and EnumCollection
-- 2.1:
--      - NOTE: This version introduces breaking changes!
--      - Replaced all APPLICATION tags for fields with CONTEXT-
SPECIFIC tags
--      APPLICATION tags are only used for custom types now.
-- 2.0:
--      Initial Release
--
```

```

EmberPlus-Glow DEFINITIONS EXPLICIT TAGS ::= BEGIN

-- =====
--
-- Primitive Types
--
-- =====

EmberString ::= UTF8String
Integer32 ::= INTEGER (-2147483648 .. 2147483647)
Integer64 ::= INTEGER (-9223372036854775808 ..
9223372036854775807)

-- this is the base oid for all RELATIVE-OID values defined
-- in this document.
-- when using the RELATIVE-OID type, defining a base oid is
-- required by ASN.1.
-- does not have any impact upon the DTD.
baseOid OBJECT IDENTIFIER ::= { iso(1) org(3) dod(6)
internet(1) private(4) enterprises(1) lsb(37411) lsb-mgmt(2)
emberPlus(1) glow(1) glowVolatile(100) }

-- =====
--
-- Template
--
-- =====

Template ::=
    [APPLICATION 24] IMPLICIT
        SET {
            number          [0] Integer32,
            element         [1] TemplateElement,
            description      [2] EmberString OPTIONAL,

QualifiedTemplate ::=
    [APPLICATION 25] IMPLICIT SET {
        path               [0] RELATIVE-OID
        element            [1] TemplateElement
        description        [2] EmberString

```

```

    }

-- =====
--
-- Parameter
--
-- =====
    }

TemplateElement ::=
    CHOICE {
        parameter      Parameter,
        node            Node,
        matrix          Matrix,
        function         Function,
    }

-- =====
--
-- Parameter
--
-- =====

Parameter ::=
    [APPLICATION 1] IMPLICIT
    SEQUENCE {
        number      [0] Integer32,
        contents    [1] ParameterContents OPTIONAL,
        children    [2] ElementCollection OPTIONAL
    }

QualifiedParameter ::=
    [APPLICATION 9] IMPLICIT
    SEQUENCE {
        path        [0] RELATIVE-OID,
        contents    [1] ParameterContents OPTIONAL,
        children    [2] ElementCollection OPTIONAL
    }

ParameterContents ::=
    SET {
        identifier   [ 0] EmberString
OPTIONAL,

```



```

        description      [ 1] EmberString
OPTIONAL,
        value            [ 2] Value
OPTIONAL,
        minimum          [ 3] MinMax
OPTIONAL,
        maximum          [ 4] MinMax
OPTIONAL,
        access           [ 5] ParameterAccess
OPTIONAL,
        format           [ 6] EmberString
OPTIONAL,
        enumeration      [ 7] EmberString
OPTIONAL,
        factor           [ 8] Integer32
OPTIONAL,
        isOnline         [ 9] BOOLEAN
OPTIONAL,
        formula          [10] EmberString
OPTIONAL,
        step             [11] Integer32
OPTIONAL,
        default          [12] Value
OPTIONAL,
        type             [13] ParameterType
OPTIONAL,
        streamIdentifier [14] Integer32
OPTIONAL,
        enumMap          [15] StringIntegerCollection
OPTIONAL,
        streamDescriptor [16] StreamDescription
OPTIONAL,
        schemaIdentifiers[17] EmberString
OPTIONAL,
        templateReference[18] Relative OID OPTIONAL
    }

```

```

Value ::=
    CHOICE {
        integer Integer64,
        real     REAL,
        string   EmberString,
        boolean  BOOLEAN,
        octets   OCTET STRING
    }

```

```
    }

MinMax ::=
    CHOICE {
        integer Integer64,
        real     REAL
    }

ParameterType ::=
    INTEGER {
        integer (1),
        real    (2),
        string  (3),
        boolean (4),
        trigger (5),
        enum    (6),
        octets  (7)
    }

ParameterAccess ::=
    INTEGER {
        none      (0),
        read      (1), -- default
        write     (2),
        readWrite (3)
    }

StringIntegerPair ::=
    [APPLICATION 7] IMPLICIT
    SEQUENCE {
        entryString  [0] EmberString,
        entryInteger [1] Integer32
    }

StringIntegerCollection ::=
    [APPLICATION 8] IMPLICIT
    SEQUENCE OF [0] StringIntegerPair

StreamDescription ::=
```

```

[APPLICATION 12] IMPLICIT
    SEQUENCE {
        format [0] StreamFormat,
        offset [1] Integer32 -- byte offset of the value
in the streamed blob.
    }

-- type:      0=uint,  1=int,   2=float
-- size:      0=1byte, 1=2byte, 2=4byte, 3=8byte
-- endianness: 0=big,  1=little
StreamFormat ::=
    INTEGER {
        unsignedInt8           ( 0), -- 00000 00 0
        unsignedInt16BigEndian ( 2), -- 00000 01 0
        unsignedInt16LittleEndian ( 3), -- 00000 01 1
        unsignedInt32BigEndian ( 4), -- 00000 10 0
        unsignedInt32LittleEndian ( 5), -- 00000 10 1
        unsignedInt64BigEndian ( 6), -- 00000 11 0
        unsignedInt64LittleEndian ( 7), -- 00000 11 1
        signedInt8             ( 8), -- 00001 00 0
        signedInt16BigEndian   (10), -- 00001 01 0
        signedInt16LittleEndian (11), -- 00001 01 1
        signedInt32BigEndian   (12), -- 00001 10 0
        signedInt32LittleEndian (13), -- 00001 10 1
        signedInt64BigEndian   (14), -- 00001 11 0
        signedInt64LittleEndian (15), -- 00001 11 1
        ieeeFloat32BigEndian   (20), -- 00010 10 0
        ieeeFloat32LittleEndian (21), -- 00010 10 1
        ieeeFloat64BigEndian   (22), -- 00010 11 0
        ieeeFloat64LittleEndian (23), -- 00010 11 1
    }

-- =====
--
-- Command
--
-- =====

Command ::=
    [APPLICATION 2] IMPLICIT
        SEQUENCE {
            number      [0] CommandType,
            options      CHOICE {

```

```

        dirFieldMask [1] FieldFlags -- only valid if
number is getDirectory(32)
        invocation   [2] Invocation -- only valid if
number is invoke(33)
        } OPTIONAL
    }

```

```

CommandType ::=
    INTEGER {
        subscribe      (30),
        unsubscribe    (31),
        getDirectory    (32),
        invoke          (33)
    }

```

```

FieldFlags ::=
    INTEGER {
        sparse          (-2),
        all              (-1),
        default          ( 0), -- same as "all"
        identifier      ( 1),
        description      ( 2),
        tree              ( 3),
        value            ( 4),
        connections      ( 5)
    }

```

```

-- =====
--
-- Node
--
-- =====

```

```

Node ::=
    [APPLICATION 3] IMPLICIT
        SEQUENCE {
            number      [0] Integer32,
            contents     [1] NodeContents      OPTIONAL,
            children     [2] ElementCollection OPTIONAL
        }

```

```

QualifiedNode ::=
    [APPLICATION 10] IMPLICIT
        SEQUENCE {
            path          [0] RELATIVE-OID,
            contents      [1] NodeContents      OPTIONAL,
            children      [2] ElementCollection OPTIONAL
        }

NodeContents ::=
    SET {
        identifier          [0] EmberString  OPTIONAL,
        description         [1] EmberString  OPTIONAL,
        isRoot              [2] BOOLEAN      OPTIONAL,
        isOnline            [3] BOOLEAN      OPTIONAL, --
        default is true
        schemaIdentifiers[4] EmberString  OPTIONAL,
        templateReference[5] Relative OID OPTIONAL
    }

-- =====
--
-- Matrix
--
-- =====

Matrix ::=
    [APPLICATION 13] IMPLICIT
        SEQUENCE {
            number          [0] Integer32,
            contents        [1] MatrixContents      OPTIONAL,
            children        [2] ElementCollection    OPTIONAL,
            targets         [3] TargetCollection     OPTIONAL,
            sources         [4] SourceCollection     OPTIONAL,
            connections     [5] ConnectionCollection OPTIONAL
        }

MatrixContents ::=
    SET {
        identifier          [ 0] EmberString,
        description         [ 1] EmberString
OPTIONAL,
        type                [ 2] MatrixType

```

```

OPTIONAL,
    addressingMode          [ 3] MatrixAddressingMode
OPTIONAL,
    targetCount             [ 4] Integer32,
-- linear: matrix X size; nonLinear: number of targets
    sourceCount             [ 5] Integer32,
-- linear: matrix Y size; nonLinear: number of sources
    maximumTotalConnects   [ 6] Integer32
OPTIONAL, -- nToN: max number of set connections
    maximumConnectsPerTarget [ 7] Integer32
OPTIONAL, -- nToN: max number of sources connected to one
target
    parametersLocation     [ 8] ParametersLocation
OPTIONAL,
    gainParameterNumber    [ 9] Integer32
OPTIONAL, -- nToN: number of connection gain parameter
    labels                  [10] LabelCollection
OPTIONAL,
    schemaIdentifiers       [11] EmberString
OPTIONAL,
    templateReference       [12] Relative OID OPTIONAL
    }

-- Addressing scheme for node at
MatrixContents.parametersLocation:
-- N 0001 targets.<targetNumber>: subtree containing
parameters attached to target with <targetNumber>
-- N 0002 sources.<sourceNumber>: subtree containing
parameters attached to source with <targetNumber>
-- N 0003 connections.<targetNumber>.<sourceNumber>: :
subtree containing parameters attached to connection
<targetNumber>/<sourceNumber>

MatrixType ::=
    INTEGER {
        oneToN      (0),  -- default
        oneToOne    (1),
        nToN        (2)
    }

MatrixAddressingMode ::=
    INTEGER {
        linear      (0),  -- default

```

```
        nonLinear (1)
    }

ParametersLocation ::=
    CHOICE {
        basePath RELATIVE-OID, -- absolute path to node
        containing parameters for targets, sources and connections
        inline Integer32 -- subidentifier to node
        containing parameters for targets, sources and connections
    }

LabelCollection ::=
    SEQUENCE OF [0] Label

Label ::=
    [APPLICATION 18] IMPLICIT
    SEQUENCE {
        basePath [0] RELATIVE-OID,
        description [1] EmberString
    }

TargetCollection ::=
    SEQUENCE OF [0] Target

Target ::=
    [APPLICATION 14] IMPLICIT
    Signal

Signal ::=
    SEQUENCE {
        number [0] Integer32
    }

SourceCollection ::=
    SEQUENCE OF [0] Source

Source ::=
```

```

[APPLICATION 15] IMPLICIT
    Signal

ConnectionCollection ::=
    SEQUENCE OF [0] Connection

Connection ::=
    [APPLICATION 16] IMPLICIT
        SEQUENCE {
            target          [0] Integer32,
            sources          [1] PackedNumbers
OPTIONAL, -- not present or empty array means "none"
            operation       [2] ConnectionOperation
OPTIONAL,
            disposition     [3] ConnectionDisposition OPTIONAL
        }

-- Use case 1: Tally (Provider to consumer)
-- Connection: { target:1, sources:[5,2], operation:absolute,
disposition:tally }

-- Use case 2: Take (Consumer to provider)
-- Connection: { target:1, sources:[4],
operation:absolute|connect|disconnect }

-- Use case 3: TakeResponse (Provider to consumer)
-- Connection: { target:1, sources:[4], operation:absolute,
disposition:modified|pending|locked|... }

PackedNumbers ::=
    RELATIVE-OID

ConnectionOperation ::=
    INTEGER {
        absolute (0), -- default. sources contains
absolute information
        connect (1), -- nToN only. sources contains
sources to add to connection
        disconnect (2) -- nToN only. sources contains
sources to remove from connection
    }

```



```

ConnectionDisposition ::=
    INTEGER {
        tally      (0),  -- default
        modified   (1),  -- sources contains new current state
        pending    (2),  -- sources contains future state
        locked     (3)   -- error: target locked. sources
contains current state
        -- more tbd.
    }

QualifiedMatrix ::=
    [APPLICATION 17] IMPLICIT
        SEQUENCE {
            path          [0] RELATIVE-OID,
            contents      [1] MatrixContents      OPTIONAL,
            children      [2] ElementCollection   OPTIONAL,
            targets       [3] TargetCollection     OPTIONAL,
            sources       [4] SourceCollection    OPTIONAL,
            connections   [5] ConnectionCollection OPTIONAL
        }

-- =====
--
-- Function
--
-- =====

Function ::=
    [APPLICATION 19] IMPLICIT
        SEQUENCE {
            number        [0] Integer32,
            contents      [1] FunctionContents    OPTIONAL,
            children      [2] ElementCollection   OPTIONAL
        }

QualifiedFunction ::=
    [APPLICATION 20] IMPLICIT
        SEQUENCE {
            path          [0] RELATIVE-OID,
            contents      [1] FunctionContents    OPTIONAL,

```

```

        children      [2] ElementCollection  OPTIONAL
    }

FunctionContents ::=
    SET {
        identifier      [0] EmberString      OPTIONAL,
        description     [1] EmberString      OPTIONAL,
        arguments       [2] TupleDescription  OPTIONAL,
        result          [3] TupleDescription  OPTIONAL,
        templateReference[4] Relative OID     OPTIONAL,
    }

TupleDescription ::=
    SEQUENCE OF [0] TupleItemDescription

TupleItemDescription ::=
    [APPLICATION 21] IMPLICIT
    SEQUENCE {
        type            [0] ParameterType,
        name            [1] EmberString      OPTIONAL
    }

Invocation ::=
    [APPLICATION 22] IMPLICIT
    SEQUENCE {
        invocationId    [0] Integer32      OPTIONAL,
        arguments       [1] Tuple          OPTIONAL
    }

Tuple ::=
    SEQUENCE OF [0] Value

InvocationResult ::=
    [APPLICATION 23] IMPLICIT
    SEQUENCE {
        invocationId    [0] Integer32,
        success         [1] BOOLEAN        OPTIONAL,
        result          [2] Tuple          OPTIONAL
    }

```

```

-- =====
--
-- ElementCollection
--
-- =====

ElementCollection ::=
    [APPLICATION 4] IMPLICIT
        SEQUENCE OF [0] Element

Element ::=
    CHOICE {
        parameter          Parameter,
        node                Node,
        command             Command,
        matrix              Matrix,
        function            Function,
        template            Template
    }

-- =====
--
-- Streams
--
-- =====

StreamEntry ::=
    [APPLICATION 5] IMPLICIT
        SEQUENCE {
            streamIdentifier [0] Integer32,
            streamValue      [1] Value
        }

StreamCollection ::=
    [APPLICATION 6] IMPLICIT
        SEQUENCE OF [0] StreamEntry

-- =====
--

```

```
-- Root
--
-- =====

Root ::=
  [APPLICATION 0]
    CHOICE {
      elements          RootElementCollection,
      streams            StreamCollection,
      invocationResult  InvocationResult
    }

RootElementCollection ::=
  [APPLICATION 11] IMPLICIT
    SEQUENCE OF [0] RootElement

RootElement ::=
  CHOICE {
    element              Element,
    qualifiedParameter   QualifiedParameter,
    qualifiedNode         QualifiedNode,
    qualifiedMatrix       QualifiedMatrix,
    qualifiedFunction     QualifiedFunction
    qualifiedTemplate     QualifiedTemplate
  }

END
```

Appendix

S101 CRC Table

0x0000,	0x1189,	0x2312,	0x329b,	0x4624,	0x57ad,	0x6536,	0x74bf,
0x8c48,	0x9dc1,	0xaf5a,	0xbcd3,	0xca6c,	0xdb5e,	0xe97e,	0xf8f7,
0x1081,	0x0108,	0x3393,	0x221a,	0x56a5,	0x472c,	0x75b7,	0x643e,
0x9cc9,	0x8d40,	0xbfdb,	0xae52,	0xdaed,	0xcb64,	0xf9ff,	0xe876,
0x2102,	0x308b,	0x0210,	0x1399,	0x6726,	0x76af,	0x4434,	0x55bd,
0xad4a,	0xbcc3,	0x8e58,	0x9fd1,	0xeb6e,	0xfae7,	0xc87c,	0xd9f5,
0x3183,	0x200a,	0x1291,	0x0318,	0x77a7,	0x662e,	0x54b5,	0x453c,
0xbdc3,	0xac42,	0x9ed9,	0x8f50,	0xfbef,	0xea66,	0xd8fd,	0xc974,
0x4204,	0x538d,	0x6116,	0x709f,	0x0420,	0x15a9,	0x2732,	0x36bb,
0xce4c,	0xdfc5,	0xed5e,	0xfcd7,	0x8868,	0x99e1,	0xab7a,	0xbaf3,
0x5285,	0x430c,	0x7197,	0x601e,	0x14a1,	0x0528,	0x37b3,	0x263a,
0xdec3,	0xcf44,	0xfddf,	0xec56,	0x98e9,	0x8960,	0xbbfb,	0xaa72,
0x6306,	0x728f,	0x4014,	0x519d,	0x2522,	0x34ab,	0x0630,	0x17b9,
0xef4e,	0xfec7,	0xcc5c,	0xdd5,	0xa96a,	0xb8e3,	0x8a78,	0x9bf1,
0x7387,	0x620e,	0x5095,	0x411c,	0x35a3,	0x242a,	0x16b1,	0x0738,
0xffcf,	0xee46,	0xdcdd,	0xcd54,	0xb9eb,	0xa862,	0x9af9,	0x8b70,
0x8408,	0x9581,	0xa71a,	0xb693,	0xc22c,	0xd3a5,	0xe13e,	0xf0b7,
0x0840,	0x19c9,	0x2b52,	0x3adb,	0x4e64,	0x5fed,	0x6d76,	0x7cff,
0x9489,	0x8500,	0xb79b,	0xa612,	0xd2ad,	0xc324,	0xf1bf,	0xe036,
0x18c1,	0x0948,	0x3bd3,	0x2a5a,	0x5ee5,	0x4f6c,	0x7df7,	0x6c7e,
0xa50a,	0xb483,	0x8618,	0x9791,	0xe32e,	0xf2a7,	0xc03c,	0xd1b5,
0x2942,	0x38cb,	0x0a50,	0x1bd9,	0x6f66,	0x7eef,	0x4c74,	0x5dfd,
0xb58b,	0xa402,	0x9699,	0x8710,	0xf3af,	0xe226,	0xd0bd,	0xc134,
0x39c3,	0x284a,	0x1ad1,	0x0b58,	0x7fe7,	0x6e6e,	0x5cf5,	0x4d7c,
0xc60c,	0xd785,	0xe51e,	0xf497,	0x8028,	0x91a1,	0xa33a,	0xb2b3,
0x4a44,	0x5bcd,	0x6956,	0x78df,	0x0c60,	0x1de9,	0x2f72,	0x3efb,
0xd68d,	0xc704,	0xf59f,	0xe416,	0x90a9,	0x8120,	0xb3bb,	0xa232,
0x5ac5,	0x4b4c,	0x79d7,	0x685e,	0x1ce1,	0x0d68,	0x3ff3,	0x2e7a,
0xe70e,	0xf687,	0xc41c,	0xd595,	0xa12a,	0xb0a3,	0x8238,	0x93b1,
0x6b46,	0x7acf,	0x4854,	0x59dd,	0x2d62,	0x3ceb,	0x0e70,	0x1ff9,
0xf78f,	0xe606,	0xd49d,	0xc514,	0xb1ab,	0xa022,	0x92b9,	0x8330,
0x7bc7,	0x6a4e,	0x58d5,	0x495c,	0x3de3,	0x2c6a,	0x1ef1,	0x0f78

S101 Decoder Sample

```

Int  m_nRxBuffer = 0;           // Position within the Buffer
BYTE m_aRxBuffer[1024];        // The Buffer
bool m_bDataLinkEscape = false; // The Escape Flag

void Decoder::Read(BYTE Byte)
{
    if(Byte == S101_BOF)
    {
        if(m_nRxBuffer)
            TRACE("<Out of line BOF>");
    }
}

```

```

        m_nRxBuffer = 1;
        m_bDataLinkEscape = false;
        return;
    }

    if(m_nRxBuffer == 0)
        return;

    if(Byte == S101_EOF)
    {
        if(m_nRxBuffer >= 4
            && CrcCCITT16(-1, m_aRxBuffer + 1, m_nRxBuffer - 1) == 0xF0B8)
            OnMessage(m_aRxBuffer + 1, m_nRxBuffer - 3);
        else
            TRACE("<RxData Error>");

        m_nRxBuffer = 0;
        return;
    }

    if(Byte == S101_CE)
    {
        m_bDataLinkEscape = true;
        return;
    }

    if(Byte >= S101_Invalid)
        return;

    if(m_bDataLinkEscape)
    {
        m_bDataLinkEscape = false;
        Byte ^= S101_XOR;
    }

    if(m_nRxBuffer < sizeof(m_aRxBuffer))
        m_aRxBuffer[m_nRxBuffer++] = Byte;
    else
        m_nRxBuffer = 0;
}

```

S101 Encoder Sample

```

void Encoder::Write (BYTE const* pBuffer, DWORD dwLength)
{
    BYTE *pSendBuffer = new BYTE [6 + dwLength * 2];
    DWORD dwLoop;
    DWORD dwLengthNew = 0;
    pSendBuffer[dwLengthNew++] = S101_BOF;

    WORD wCRC = ~CrcCCITT16(0xFFFF, pBuffer, dwLength);

    for(dwLoop = 0; dwLoop < (int) dwLength + 2; dwLoop++)
    {
        BYTE uData;
        if(dwLoop == (int) dwLength + 0)

```

```
        uData = (BYTE) wCRC;
    else
    if(dwLoop == (int) dwLength + 1)
        uData = (BYTE) (wCRC >> 8);
    else
        uData = pBuffer[dwLoop];

    if(uData >= S101_Invalid)
    {
        pSendBuffer[dwLengthNew++] = S101_CE;
        uData ^= S101_XOR;
    }

    pSendBuffer[dwLengthNew++] = uData;
}

pSendBuffer[dwLengthNew++] = S101_EOF;
Write(pSendBuffer, dwLengthNew);
delete [] pSendBuffer;
}
```